

Hadoop Fundamentals

Hadoop

- Is an open source framework for storing, and processing extremely large data in a distributed computing environment.
- Not a database



Hadoop Components

- Distributed Computing
 - HDFS
- Parallel Processing
 - MapReduce



Advantages of Hadoop

- Low cost
 - Hadoop can run on average performing commodity
 - Help in controlling cost and achieve scalability and performance.
 - Adding or removing nodes from the cluster is simple
- Storage flexibility
 - Hadoop can store data in raw format in a distributed environment.
 - Hadoop can process the unstructured data and semi-structured data better than most of the available technologies.



Advantages of Hadoop

- Open source community
 - Hadoop is open source and supported by many contributors with a growing network of developers worldwide.
 - Many organizations such as Yahoo, Facebook, Hortonworks, and others have contributed immensely toward the progress of Hadoop and other related sub-projects.
- Fault tolerant
 - Hadoop is massively scalable and fault tolerant.
 - Hadoop is reliable in terms of data availability, and even if some nodes go down, Hadoop can recover the data.
 - Hadoop architecture assumes that nodes can go down and the system should be able to process the data.

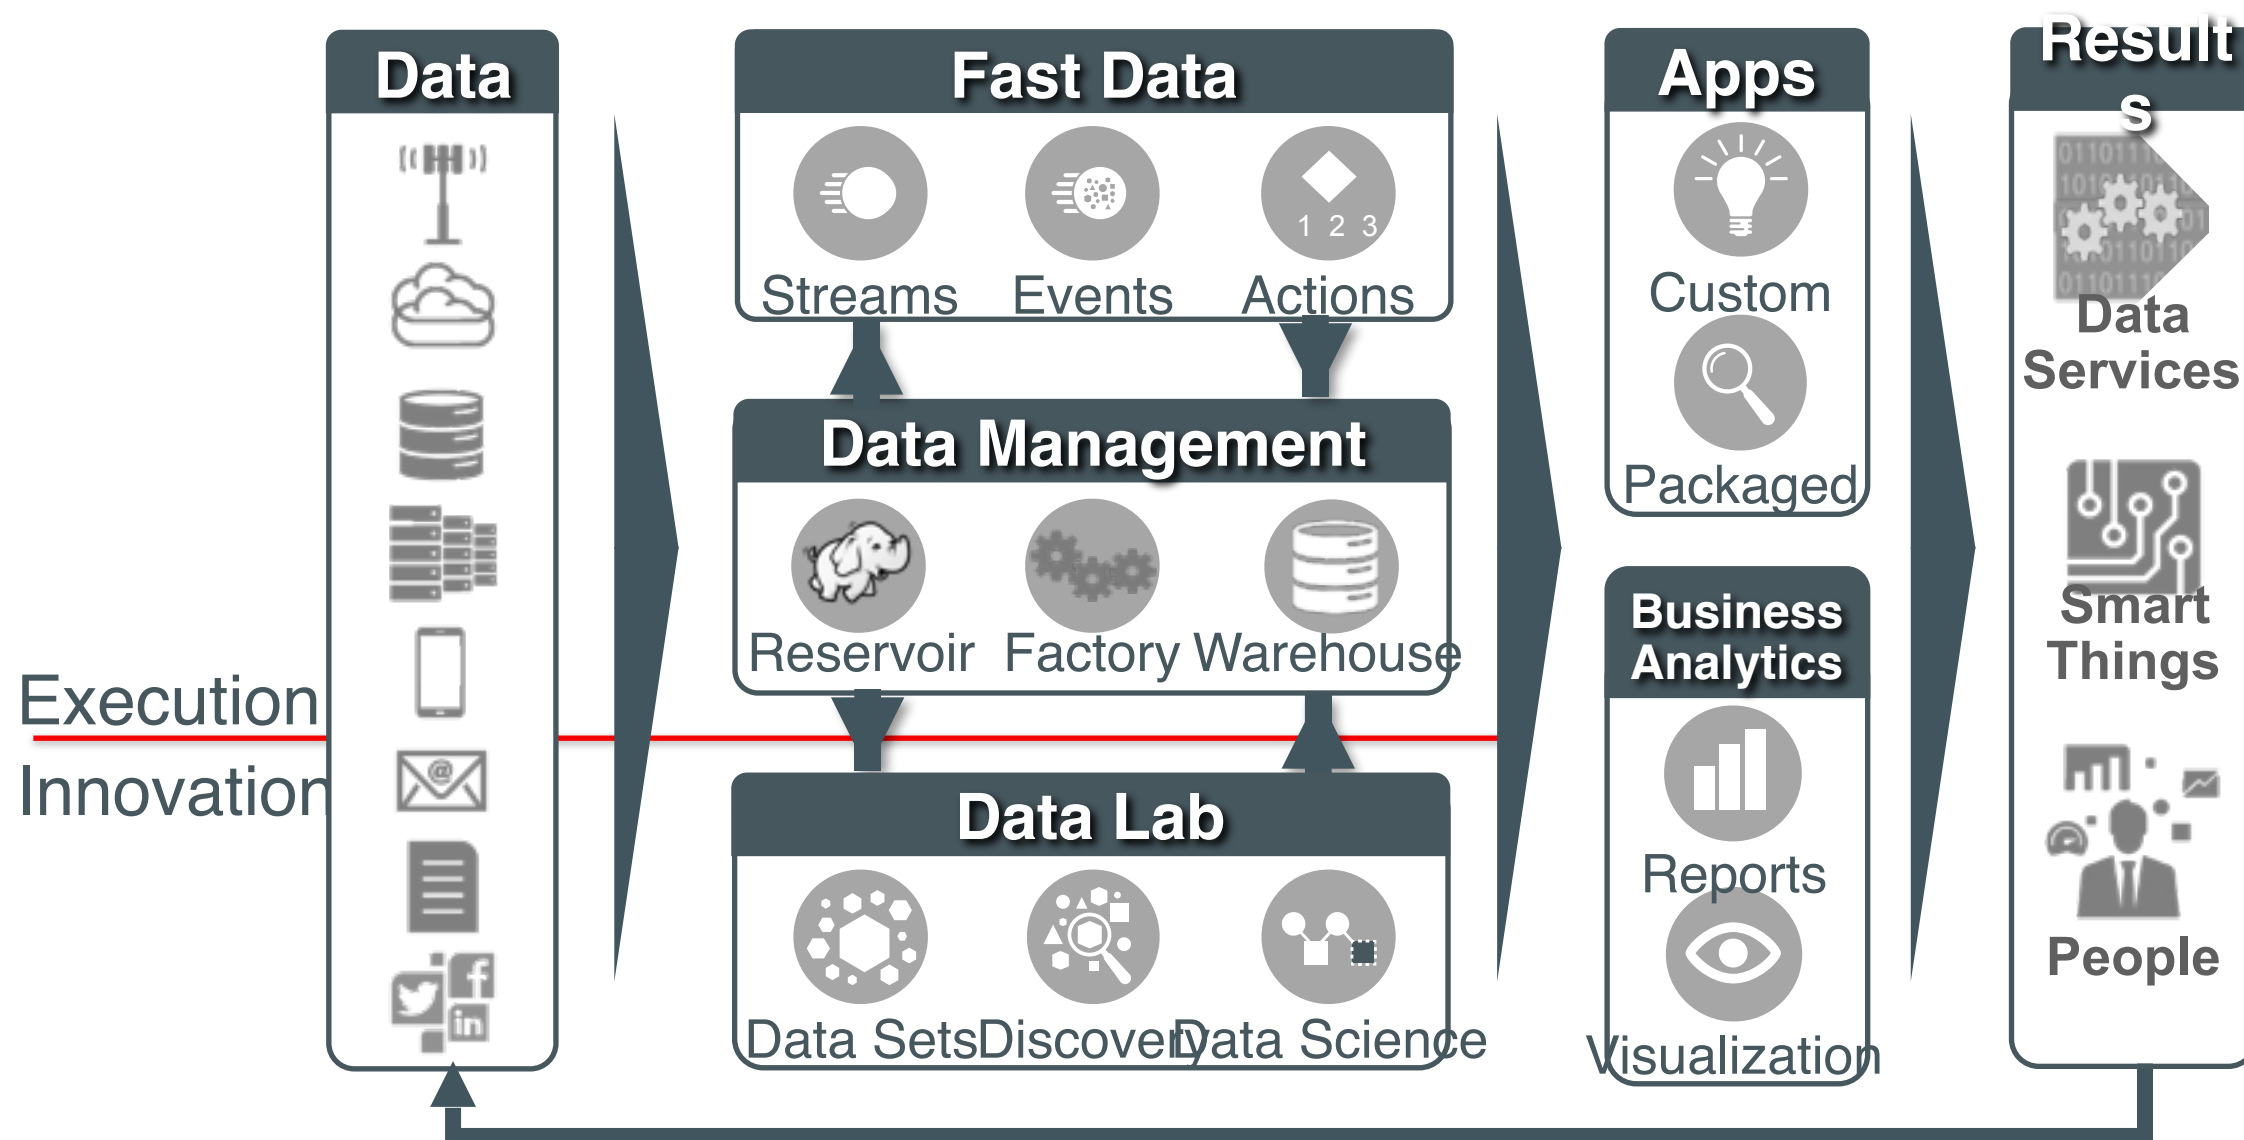


Use Cases for Hadoop

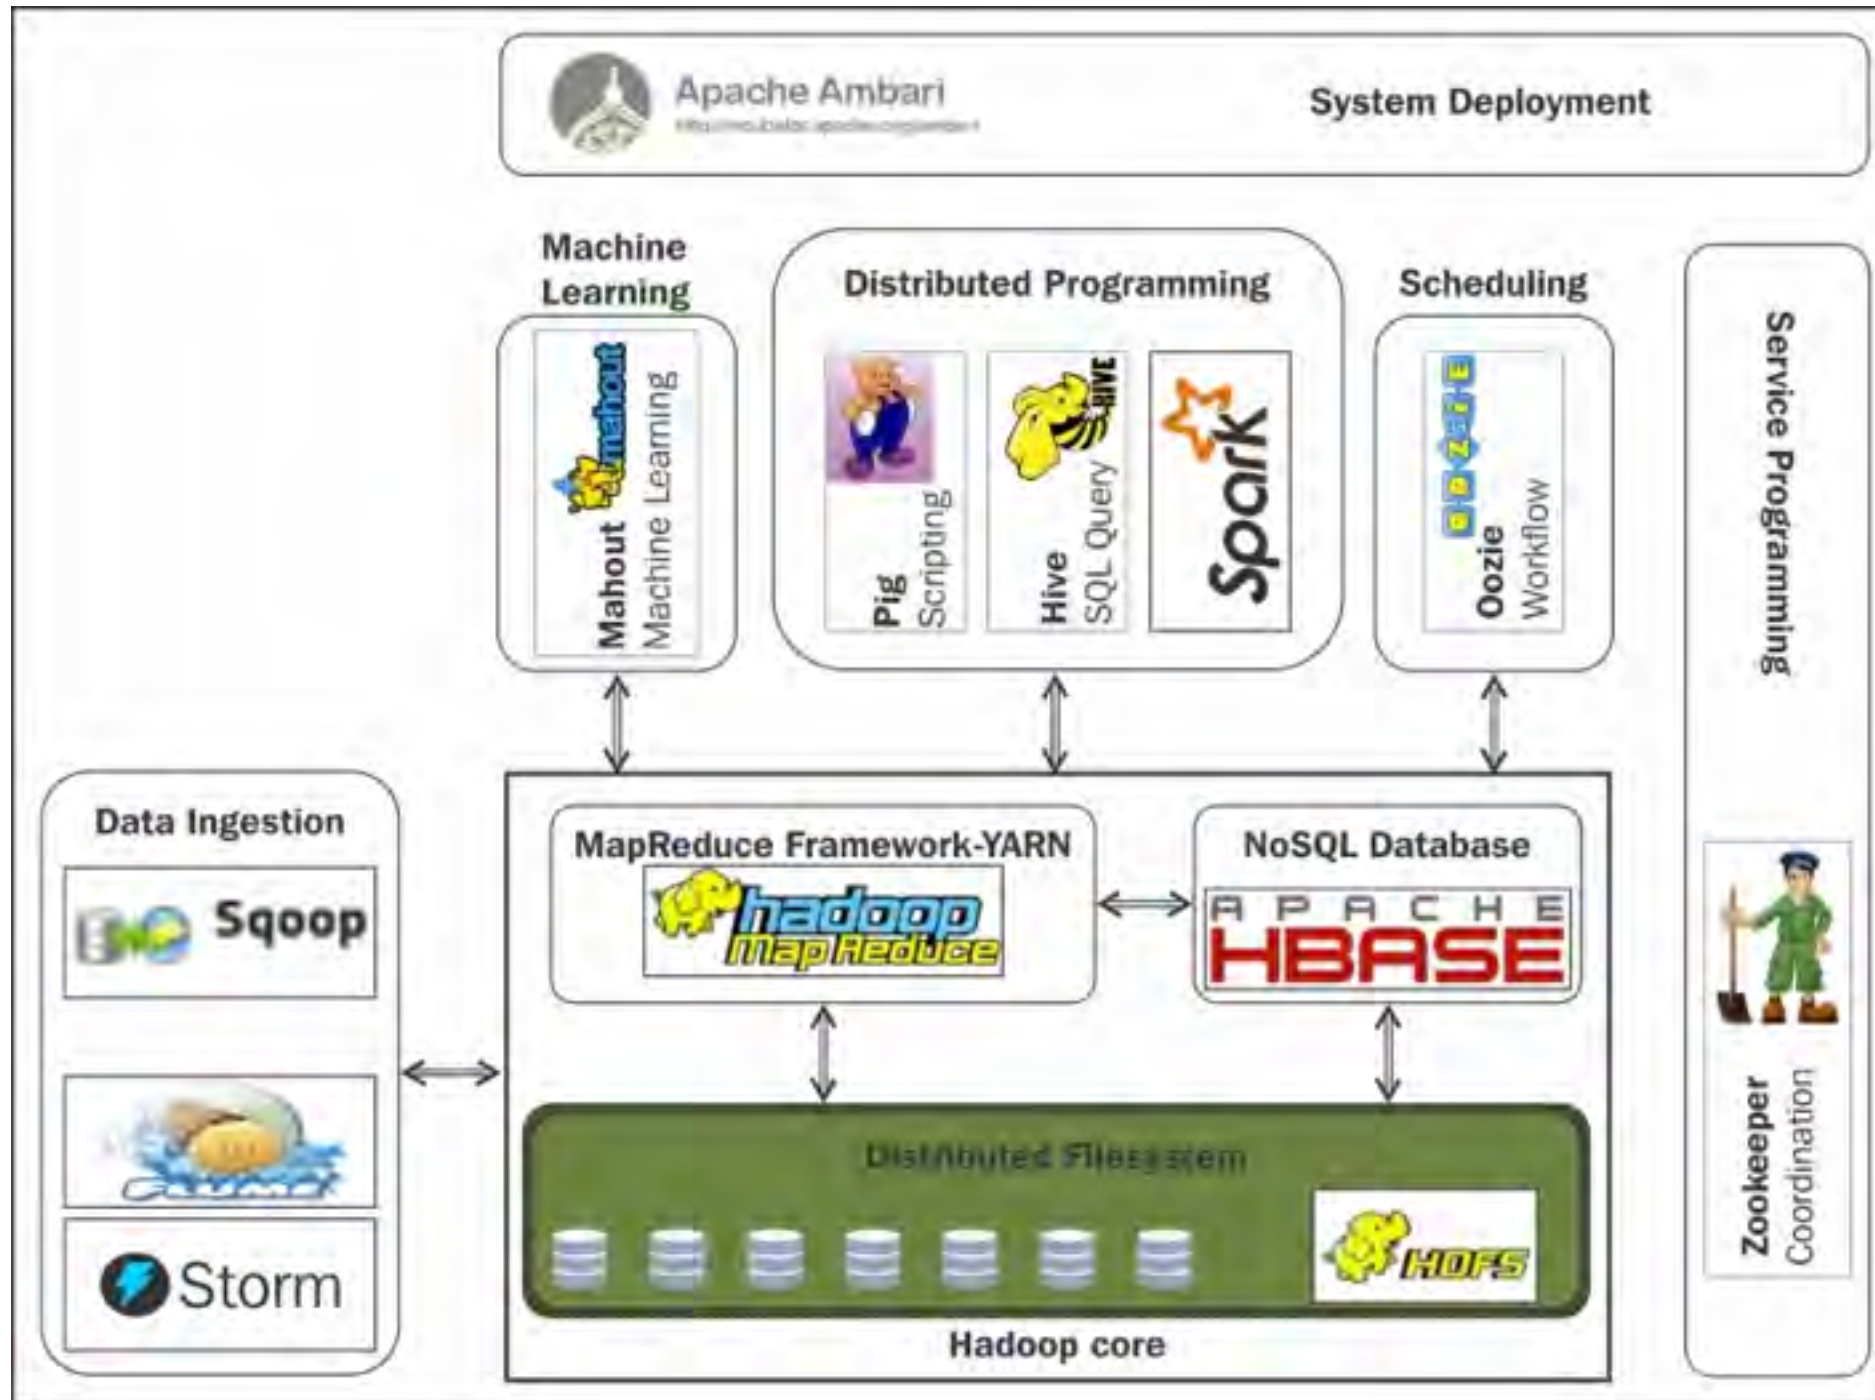
- Searching/text mining
- Log processing
- Recommendation systems
- Business intelligence/data warehousing
- Video and image analysis
- Archiving
- Graph creation and analysis
- Pattern recognition
- Risk assessment
- Sentiment analysis



Big Data Flow



Hadoop Ecosystem



- Distributed and fault tolerant file system
- Provides high throughput access to data



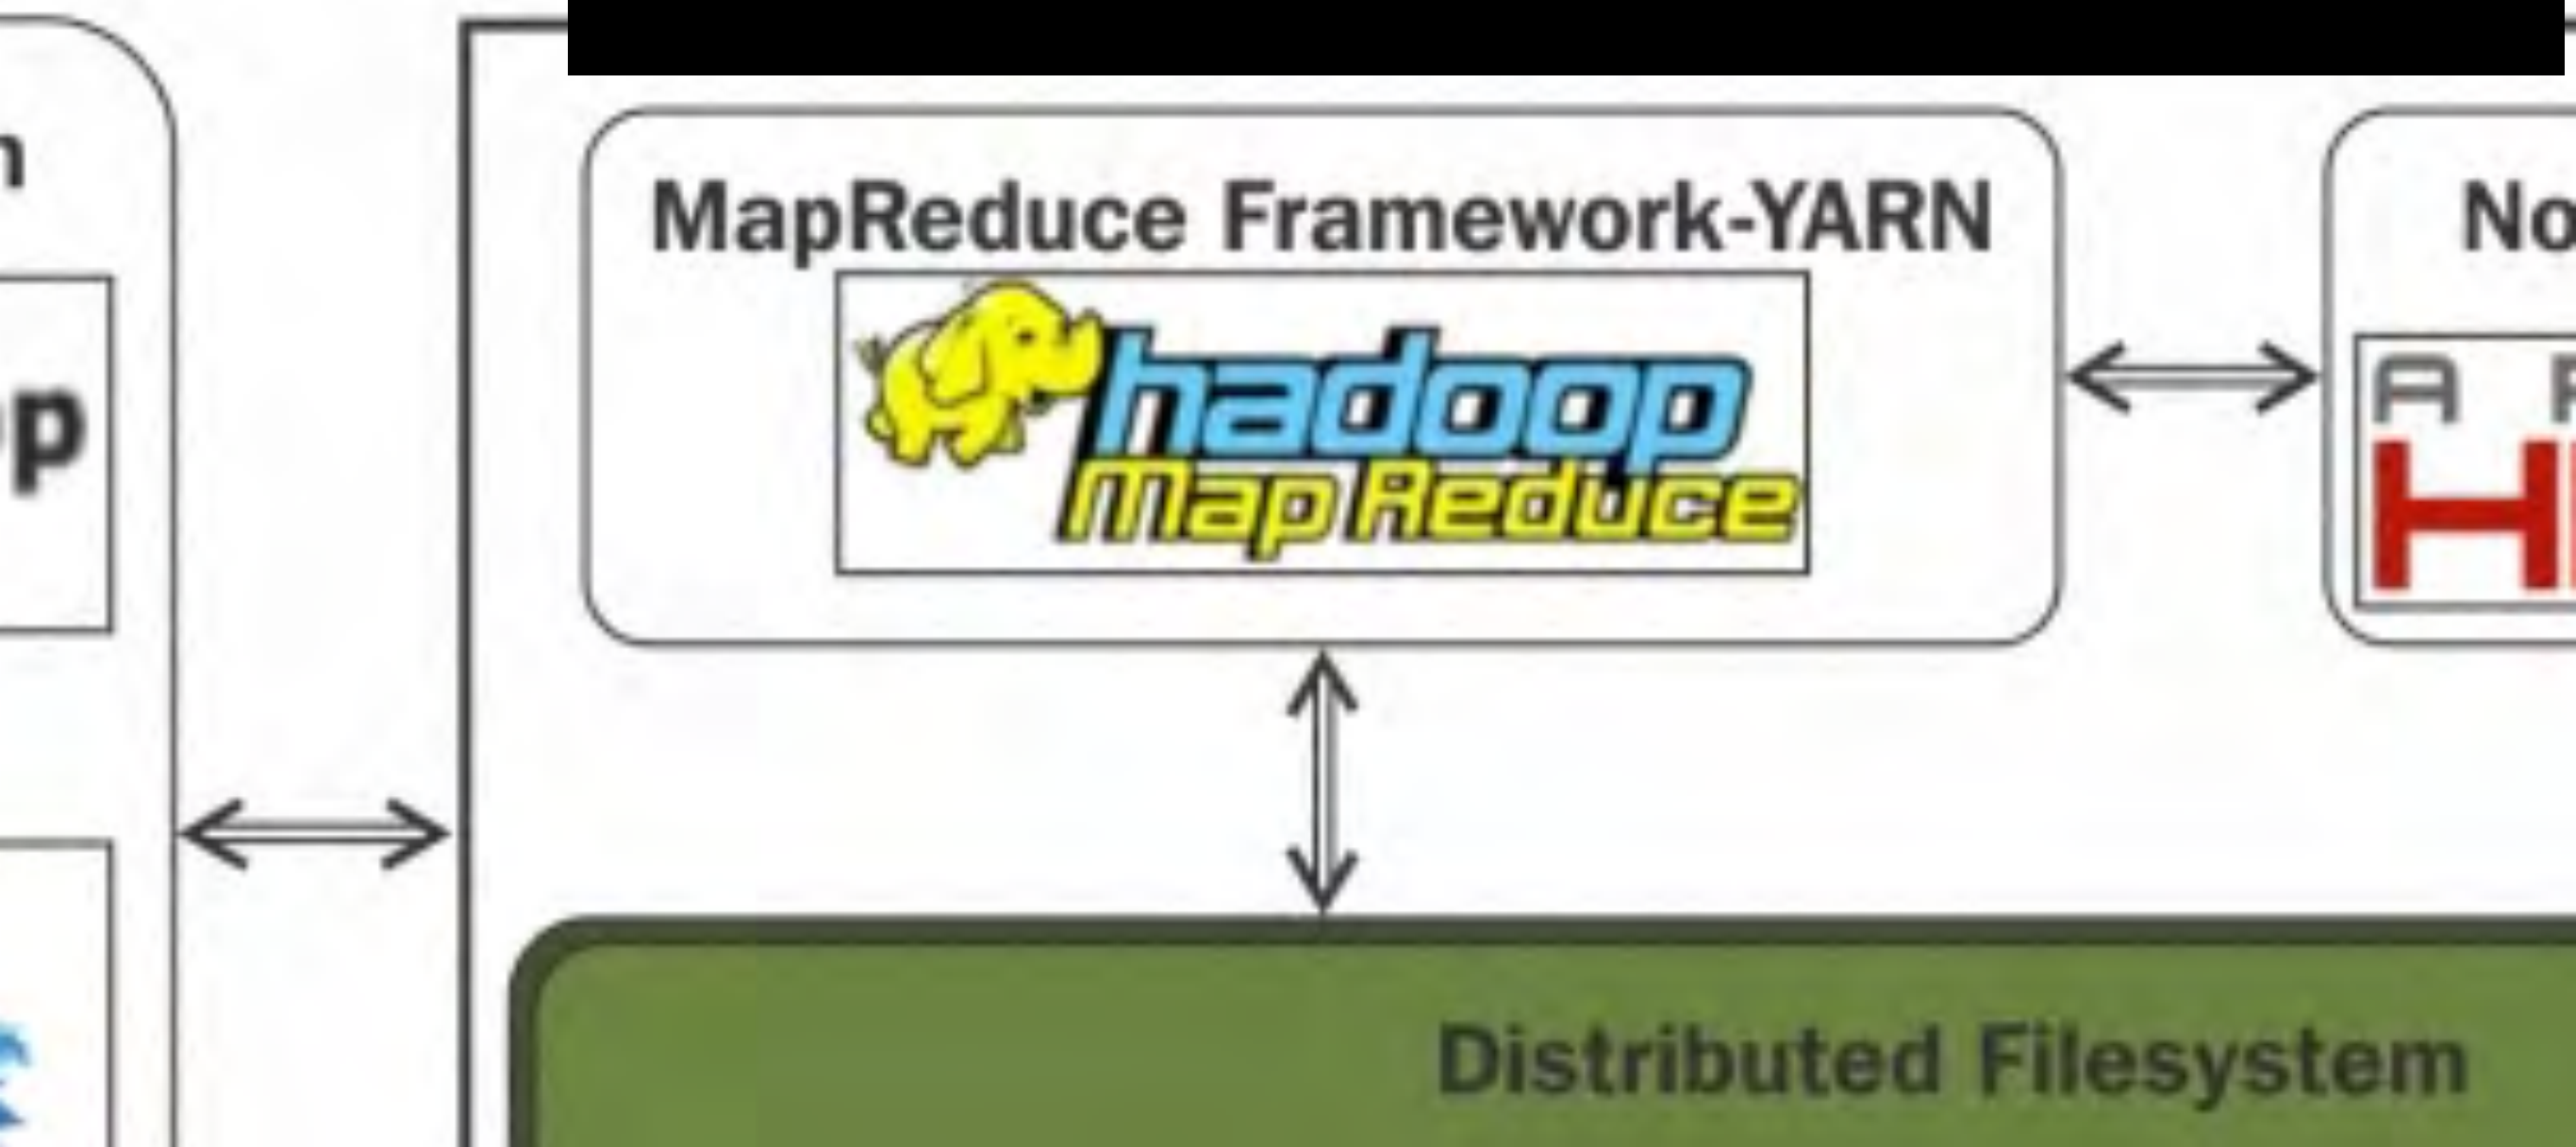
Distributed Filesystem



Hadoop core



- Parallel Processing Framework
- Processing is moved to the data



Distributed Programming



Pig
Scripting



Hive
SQL Query





- NoSQL
- Column oriented DB and a key-Value store
- Works on top of HDFS



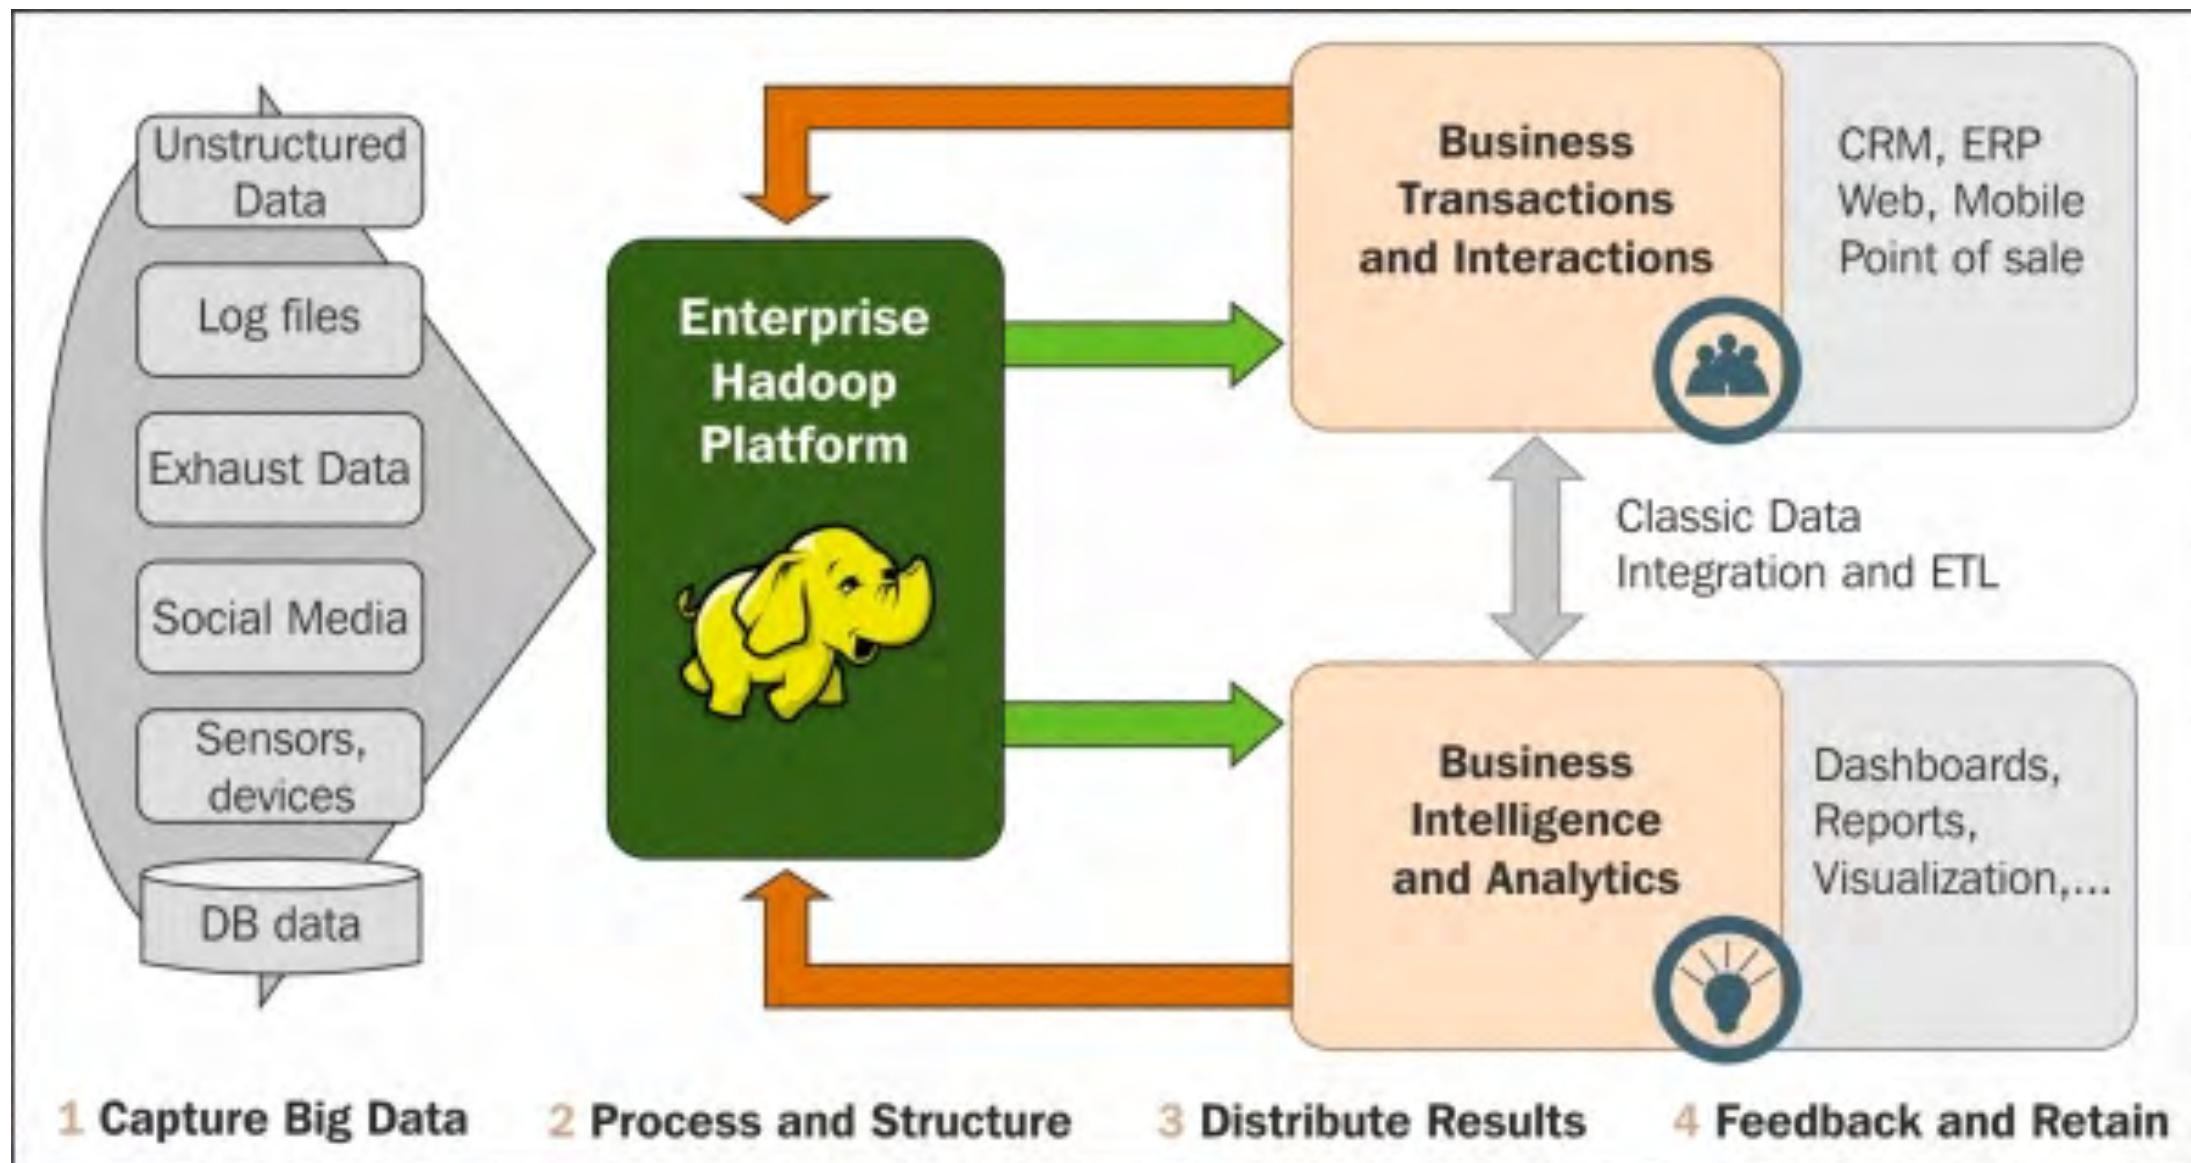
Hadoop I

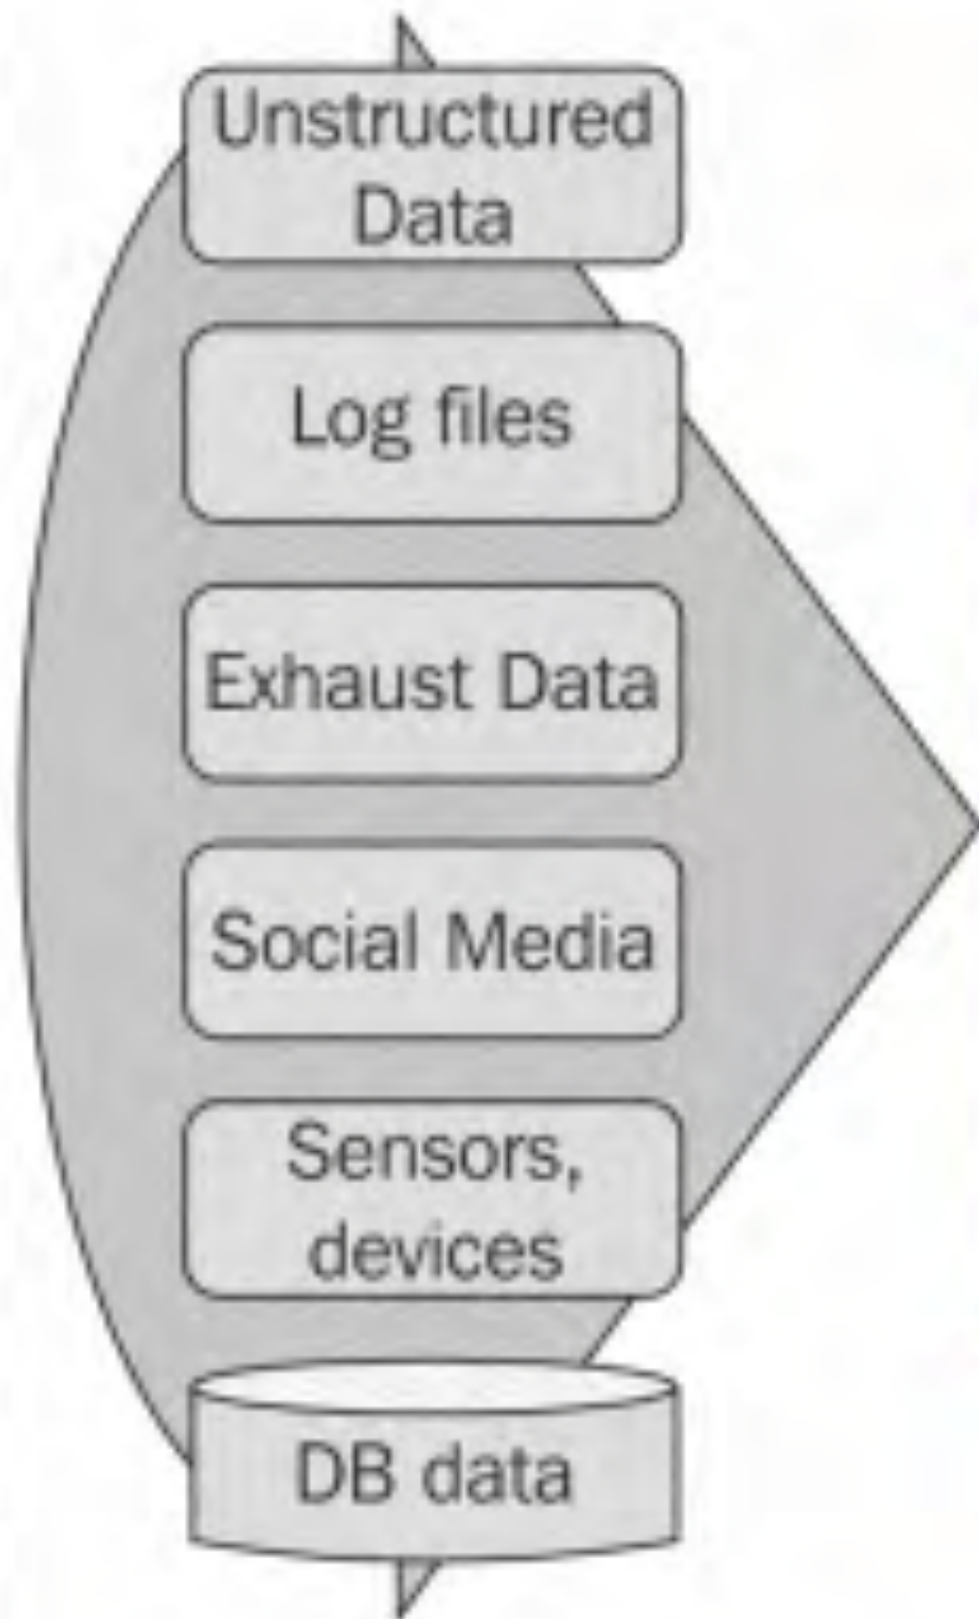
- Sqoop — Manage data between Hadoop and Relational DBs, enterprise DWs, and NoSQL
- Flume — Used for collecting, aggregating and moving log data from different sources
- Storm — Provides a solution for processing distributed data in real-time

Data Ingestion



Hadoop's data flow





1. Capture Big Data

- Sources can be extensive lists:
 - Structured, semi-structured, and unstructured,
 - some streaming, real-time data sources,
 - sensors, machine-captured data, and many other sources.

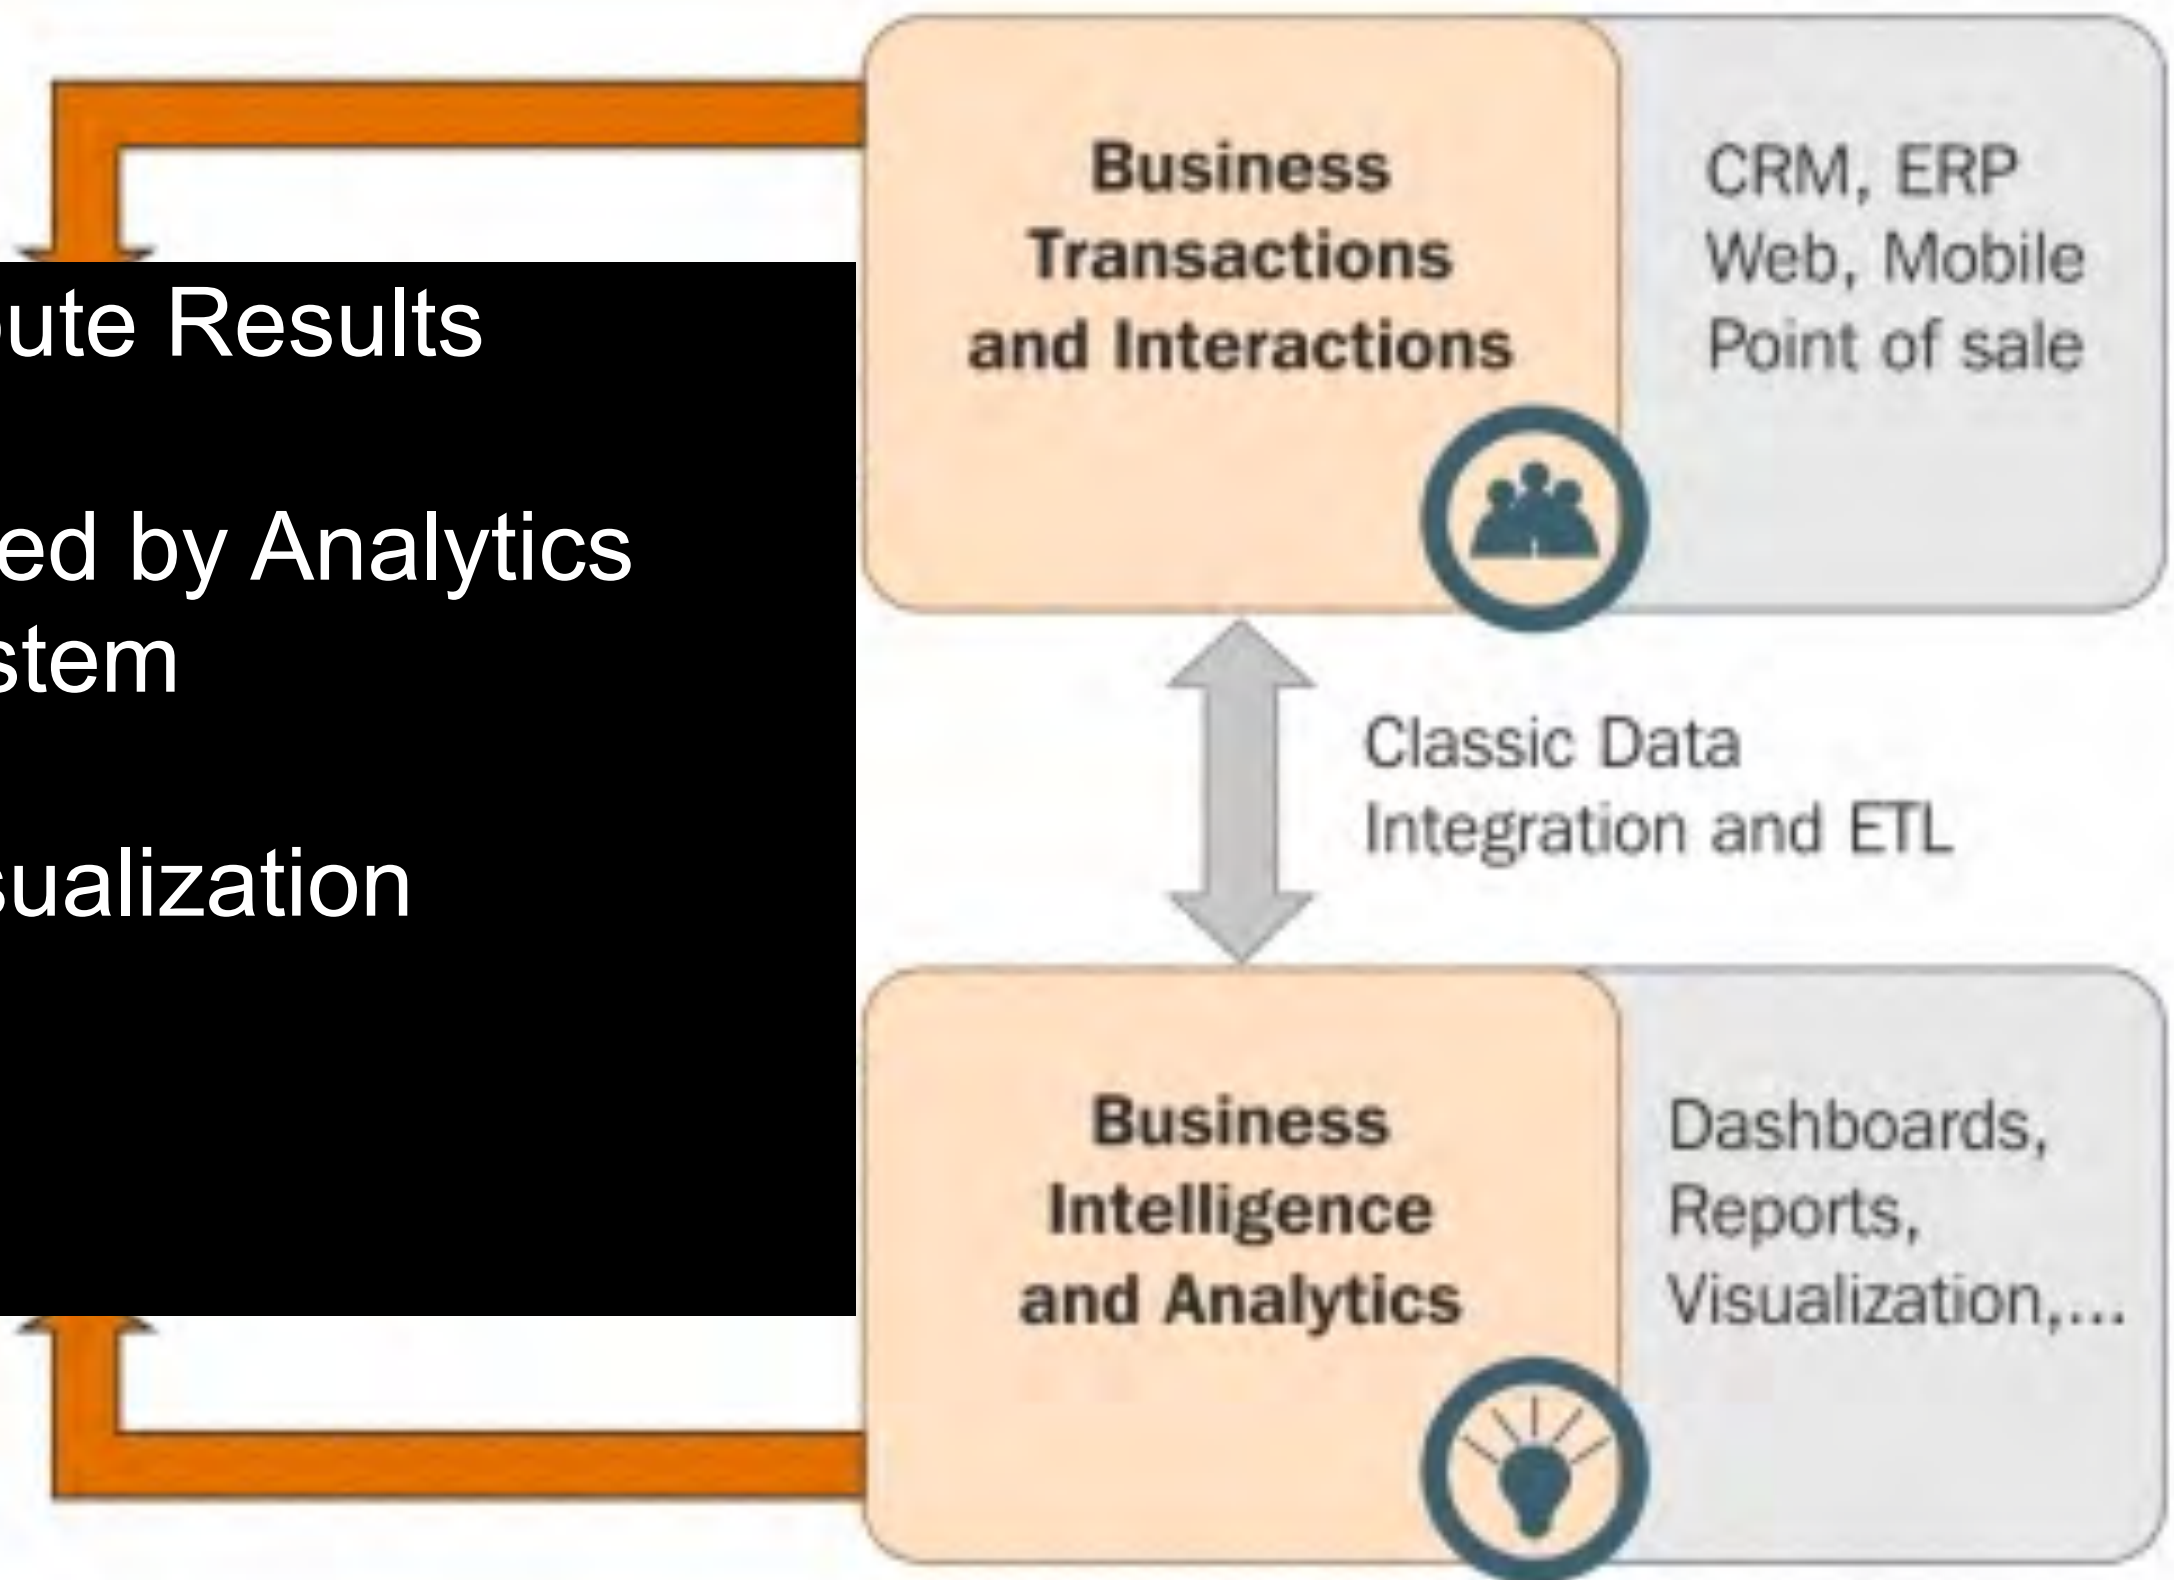
2. Process & Structure

- Cleanse
- Filter
- Transform
- Process



3. Distribute Results

- Used by Analytics system
- Visualization



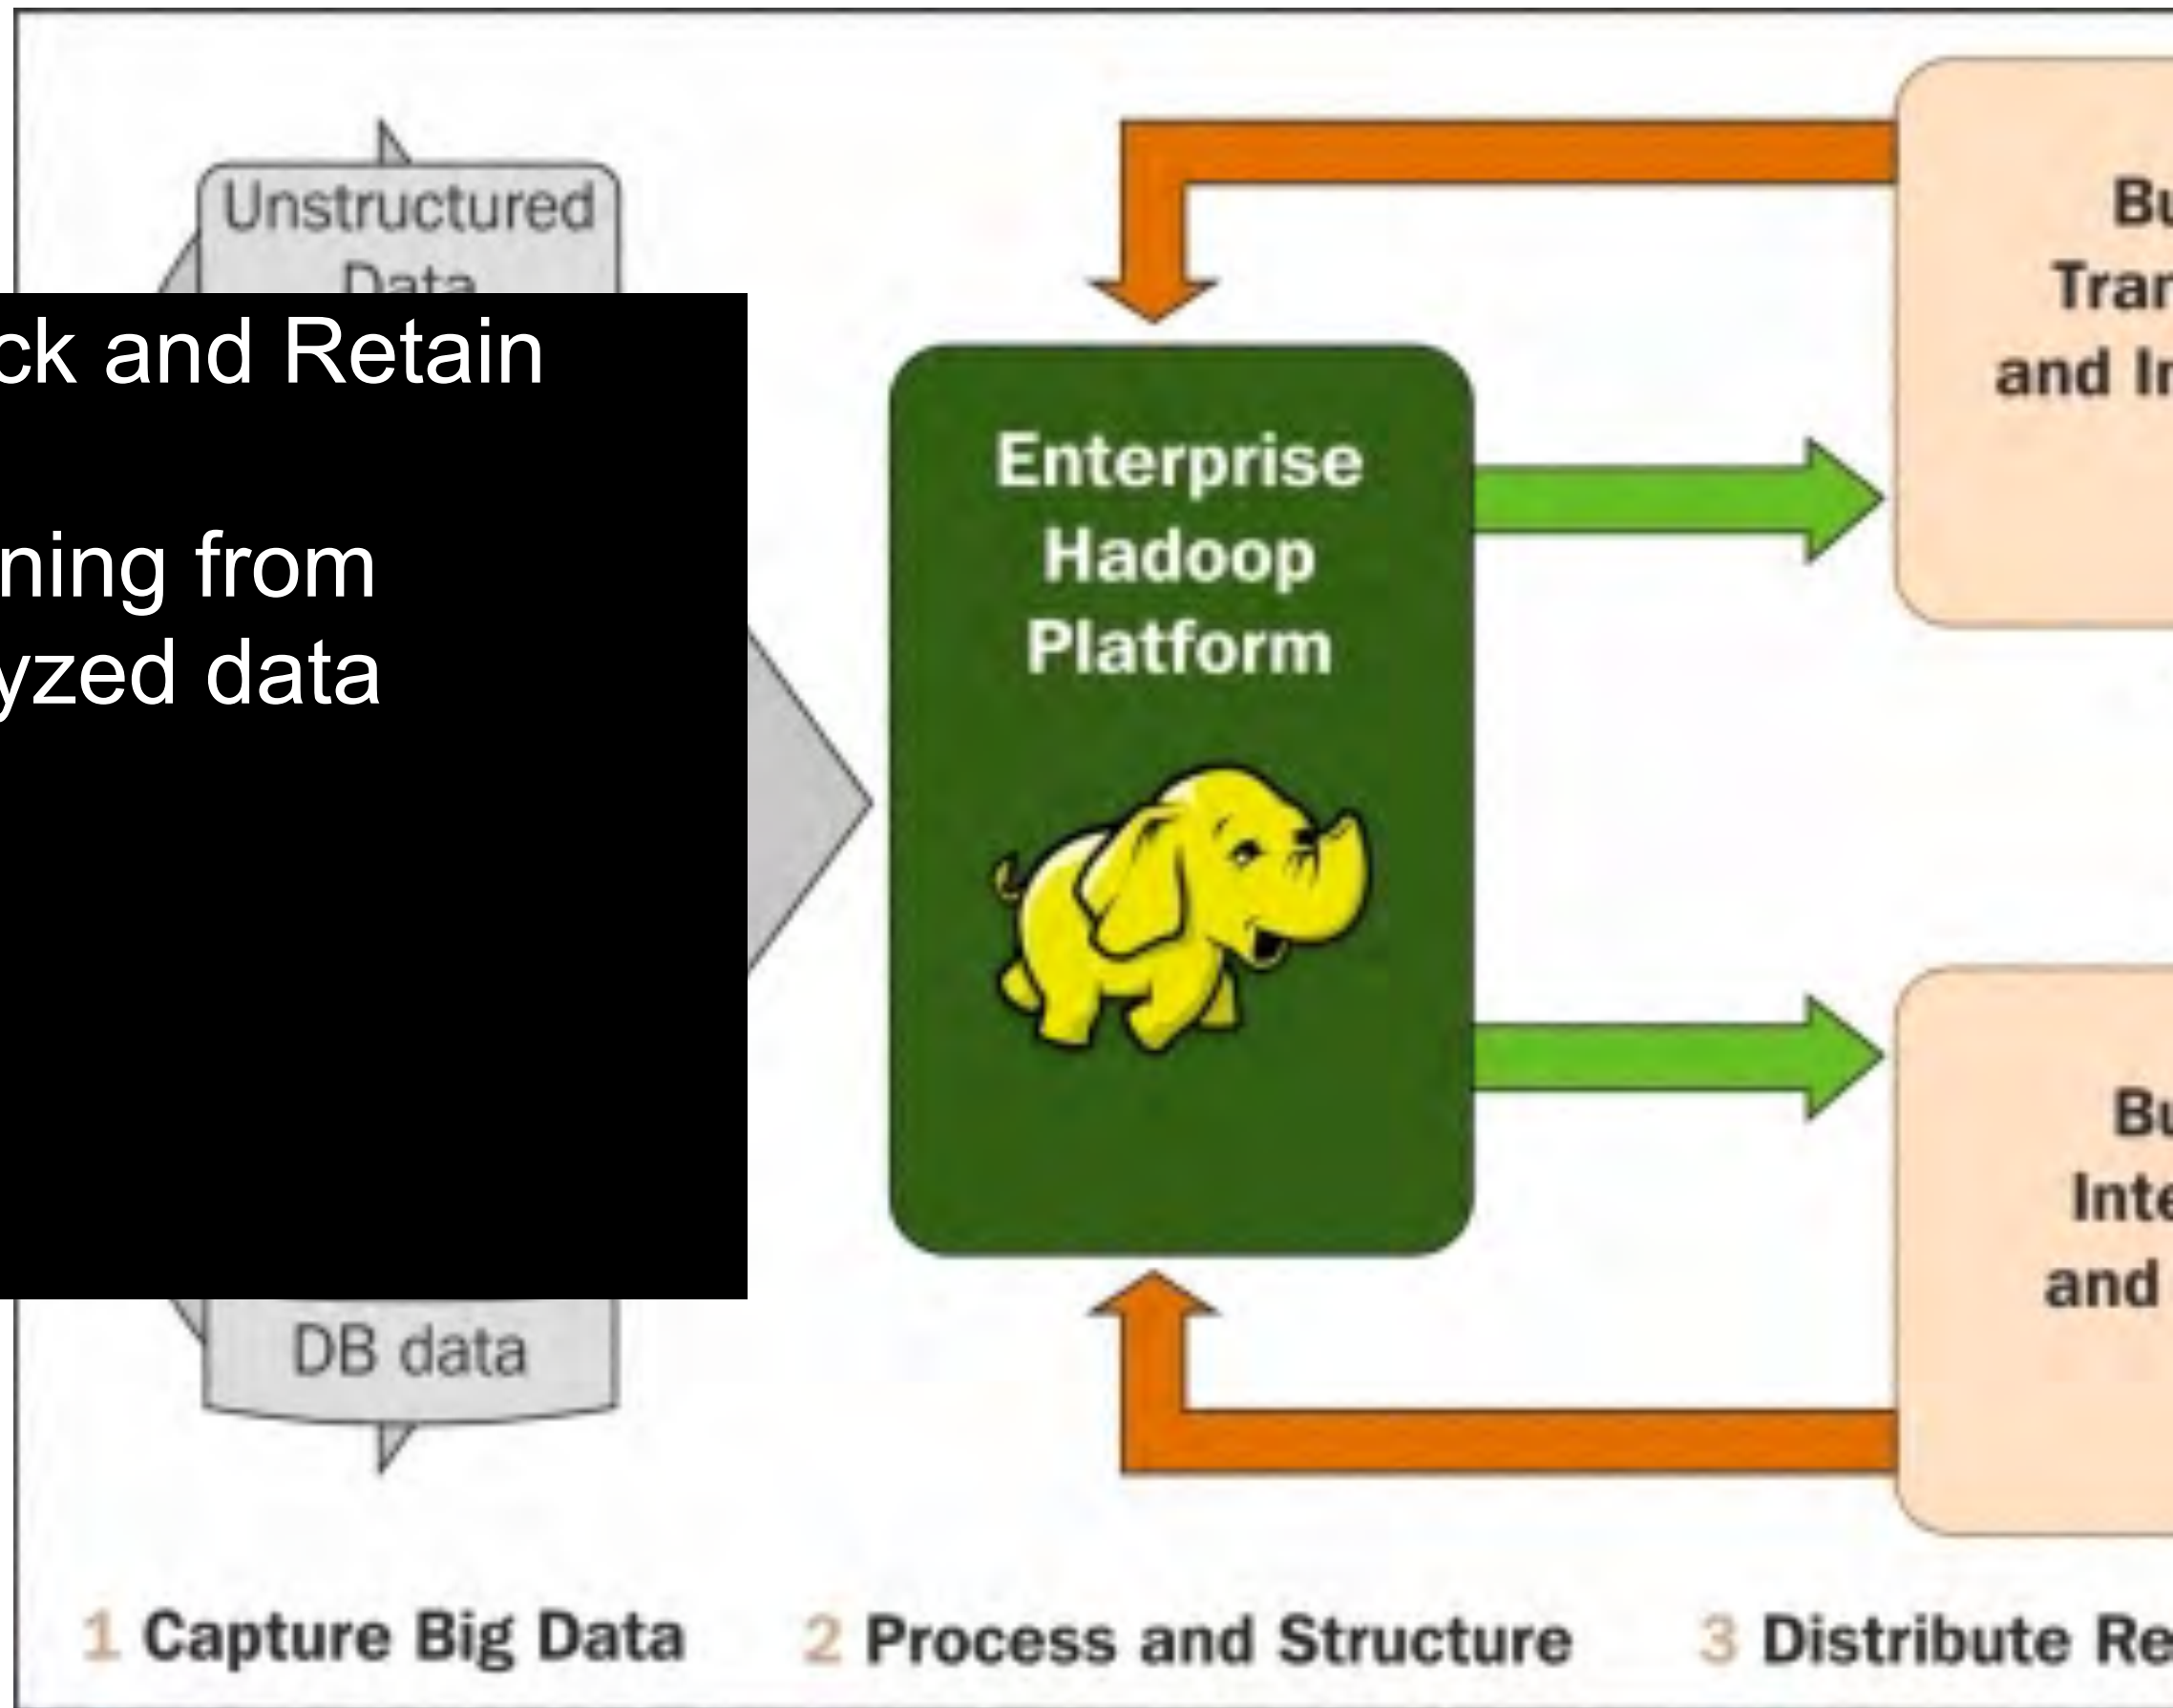
2 Process and Structure

3 Distribute Results

4 Feedback and Retain


3. Feedback and Retain

- Learning from analyzed data



HDFS

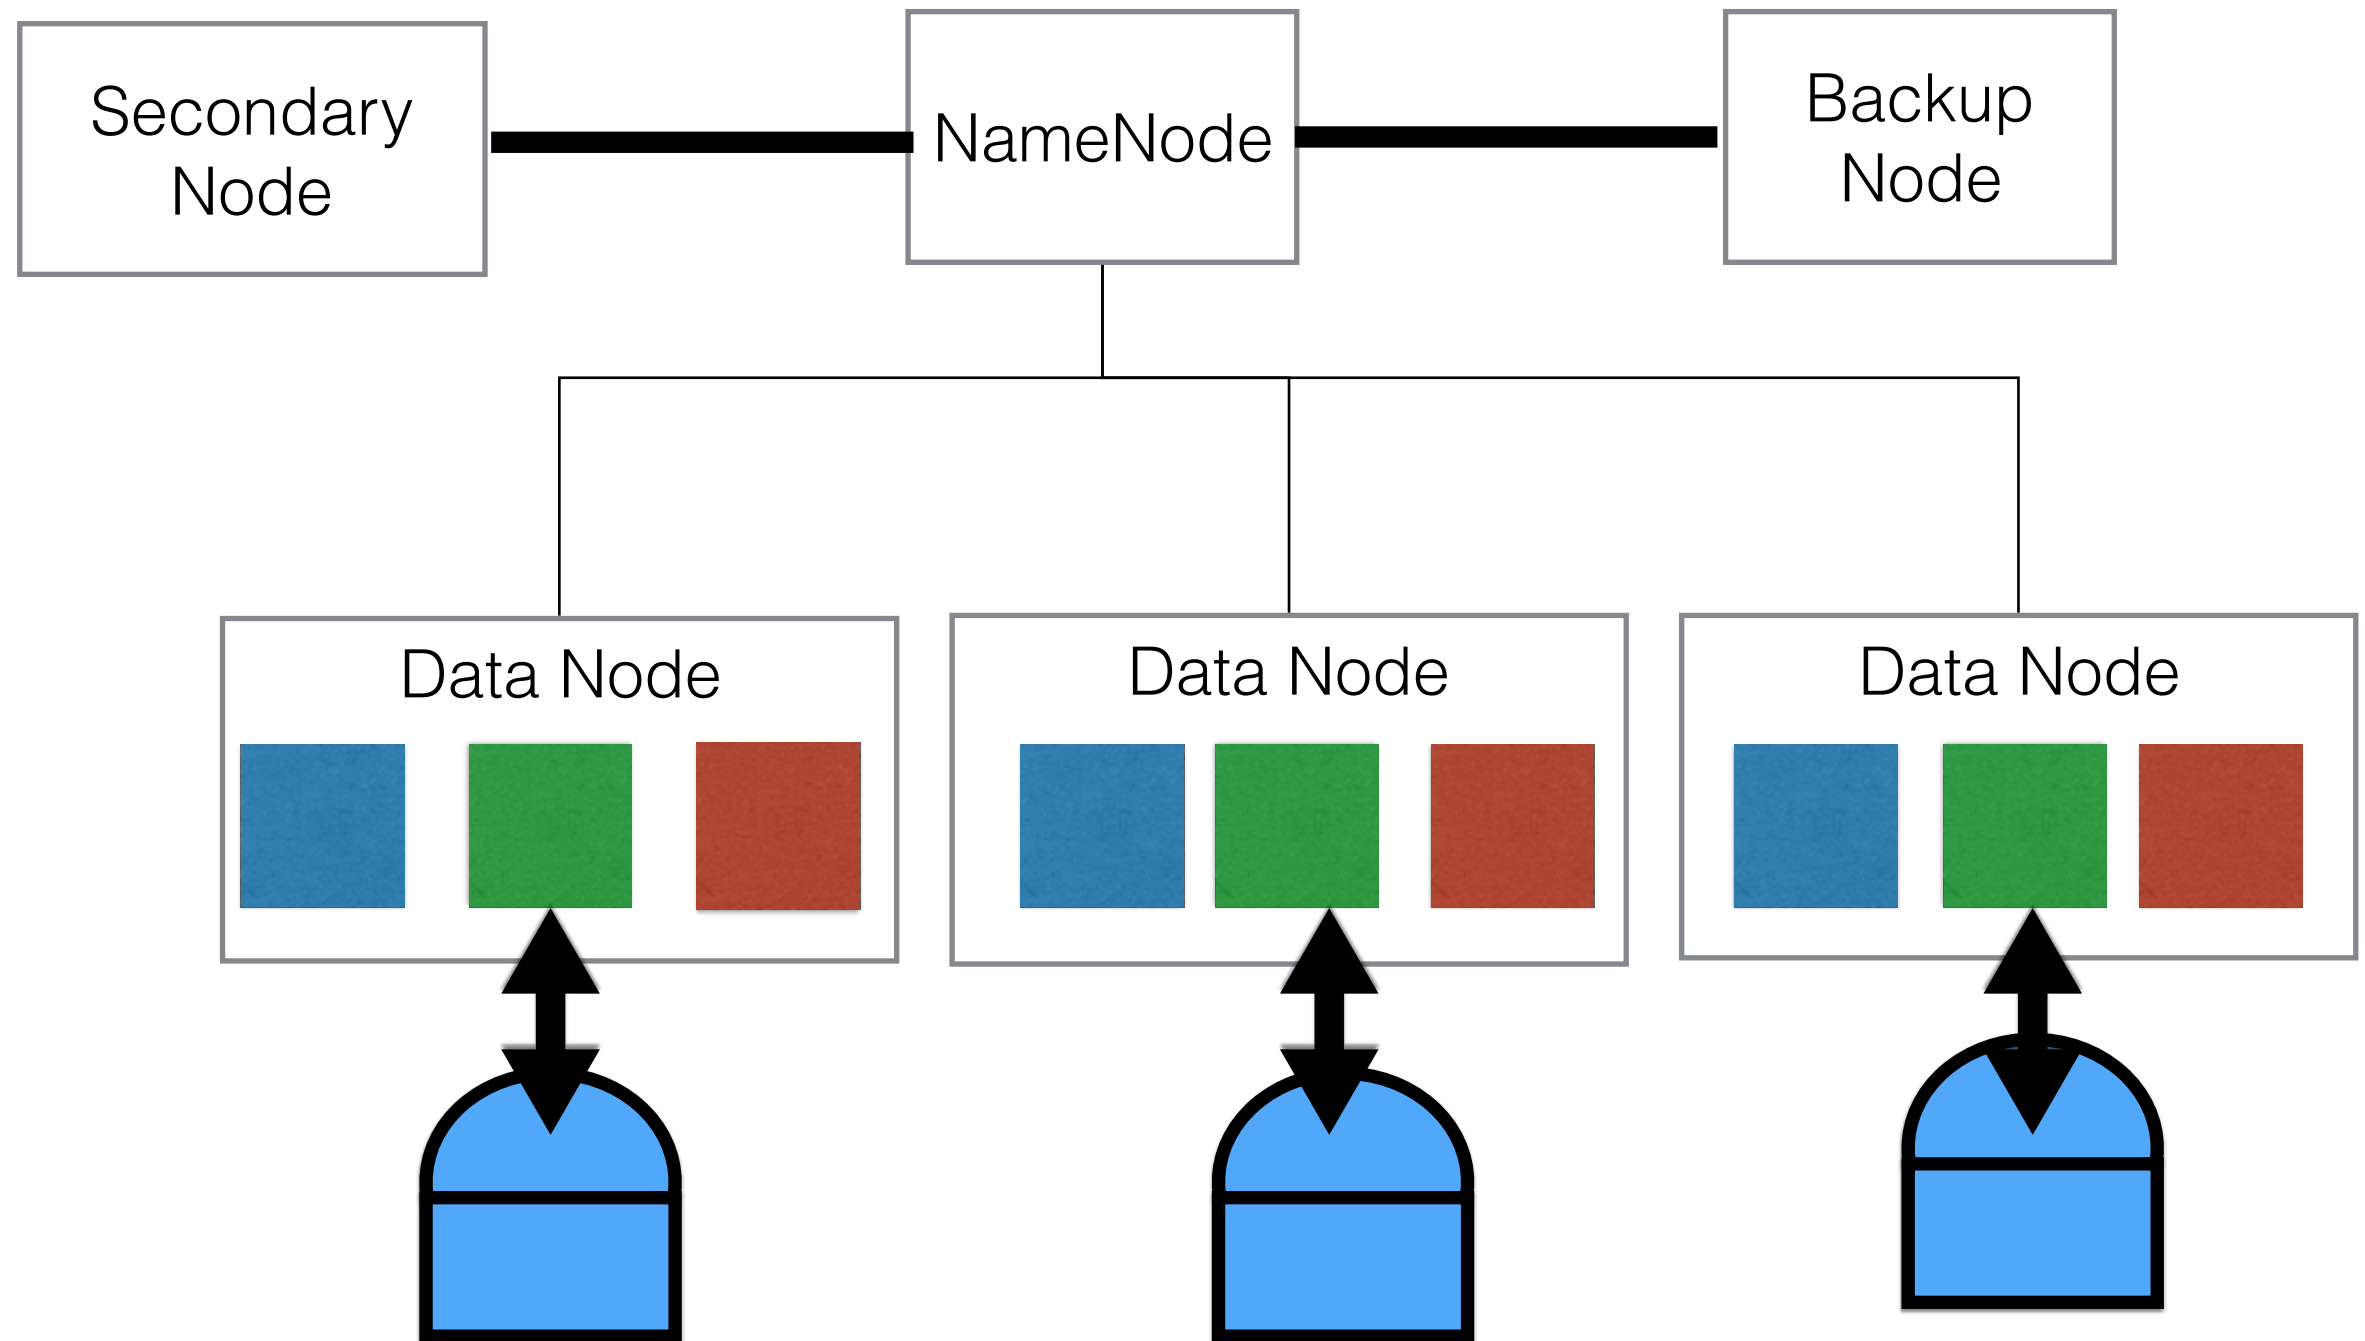
Features of HDFS

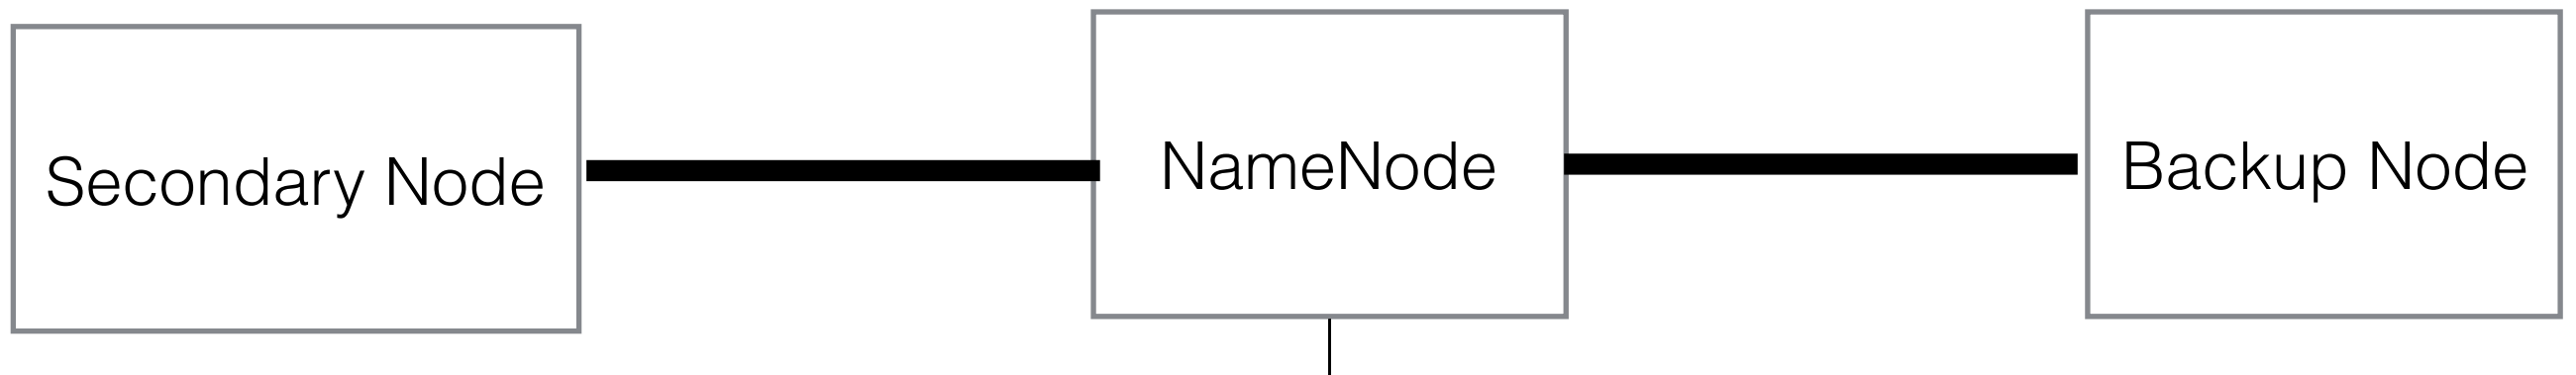


- Scalability
 - Ability to seamlessly add/delete nodes
- Reliability and fault tolerance
 - Replicates data to a configurable number of nodes
 - Data is stored in multiple nodes — system available in case of node failure
 - Replication Factor
 - ✓ Specify how many nodes you want to replicate a file
- Portability
 - Portable on different hardware and software
 - Built using the Java language — can be deployed on any machine supporting Java
- Processing closer to the data
 - Move query processing where data is rather than moving the data




HDFS Architecture





- Coordinates all storage operations including read/write in HDFS
- Manages filesystem namespace
- Holds Metadata about which nodes has what data block
 - Fimage file — Holds entire filesystem namespace including mappings of blocks to file and filesystem
 - Editlog file — Holds ever change that occurs to the filesystem metadata
- Secondary NameNode is configured for HA
 - Keeps checkpoint of FsImage and EditLog merged and available in case of a failure to NameNode
 - Requires similar configuration as the NameNode
- Backup Node
 - Similar to Secondary NameNode but keeps the FsImage in the memory
 - Always sychrnozed with the NameNode
 - Has same RAM requirements as NameNode




- Holds actual data
- Create, delete and replicate data blocks as assigned by NameNode
- DataNode sends heartbeat messages (3 sec) to the NameNode
 - No heartbeat — NameNode declares DataNode out of service
- Every 10th heartbeat is a block report
- NameNode updates metadata from block report

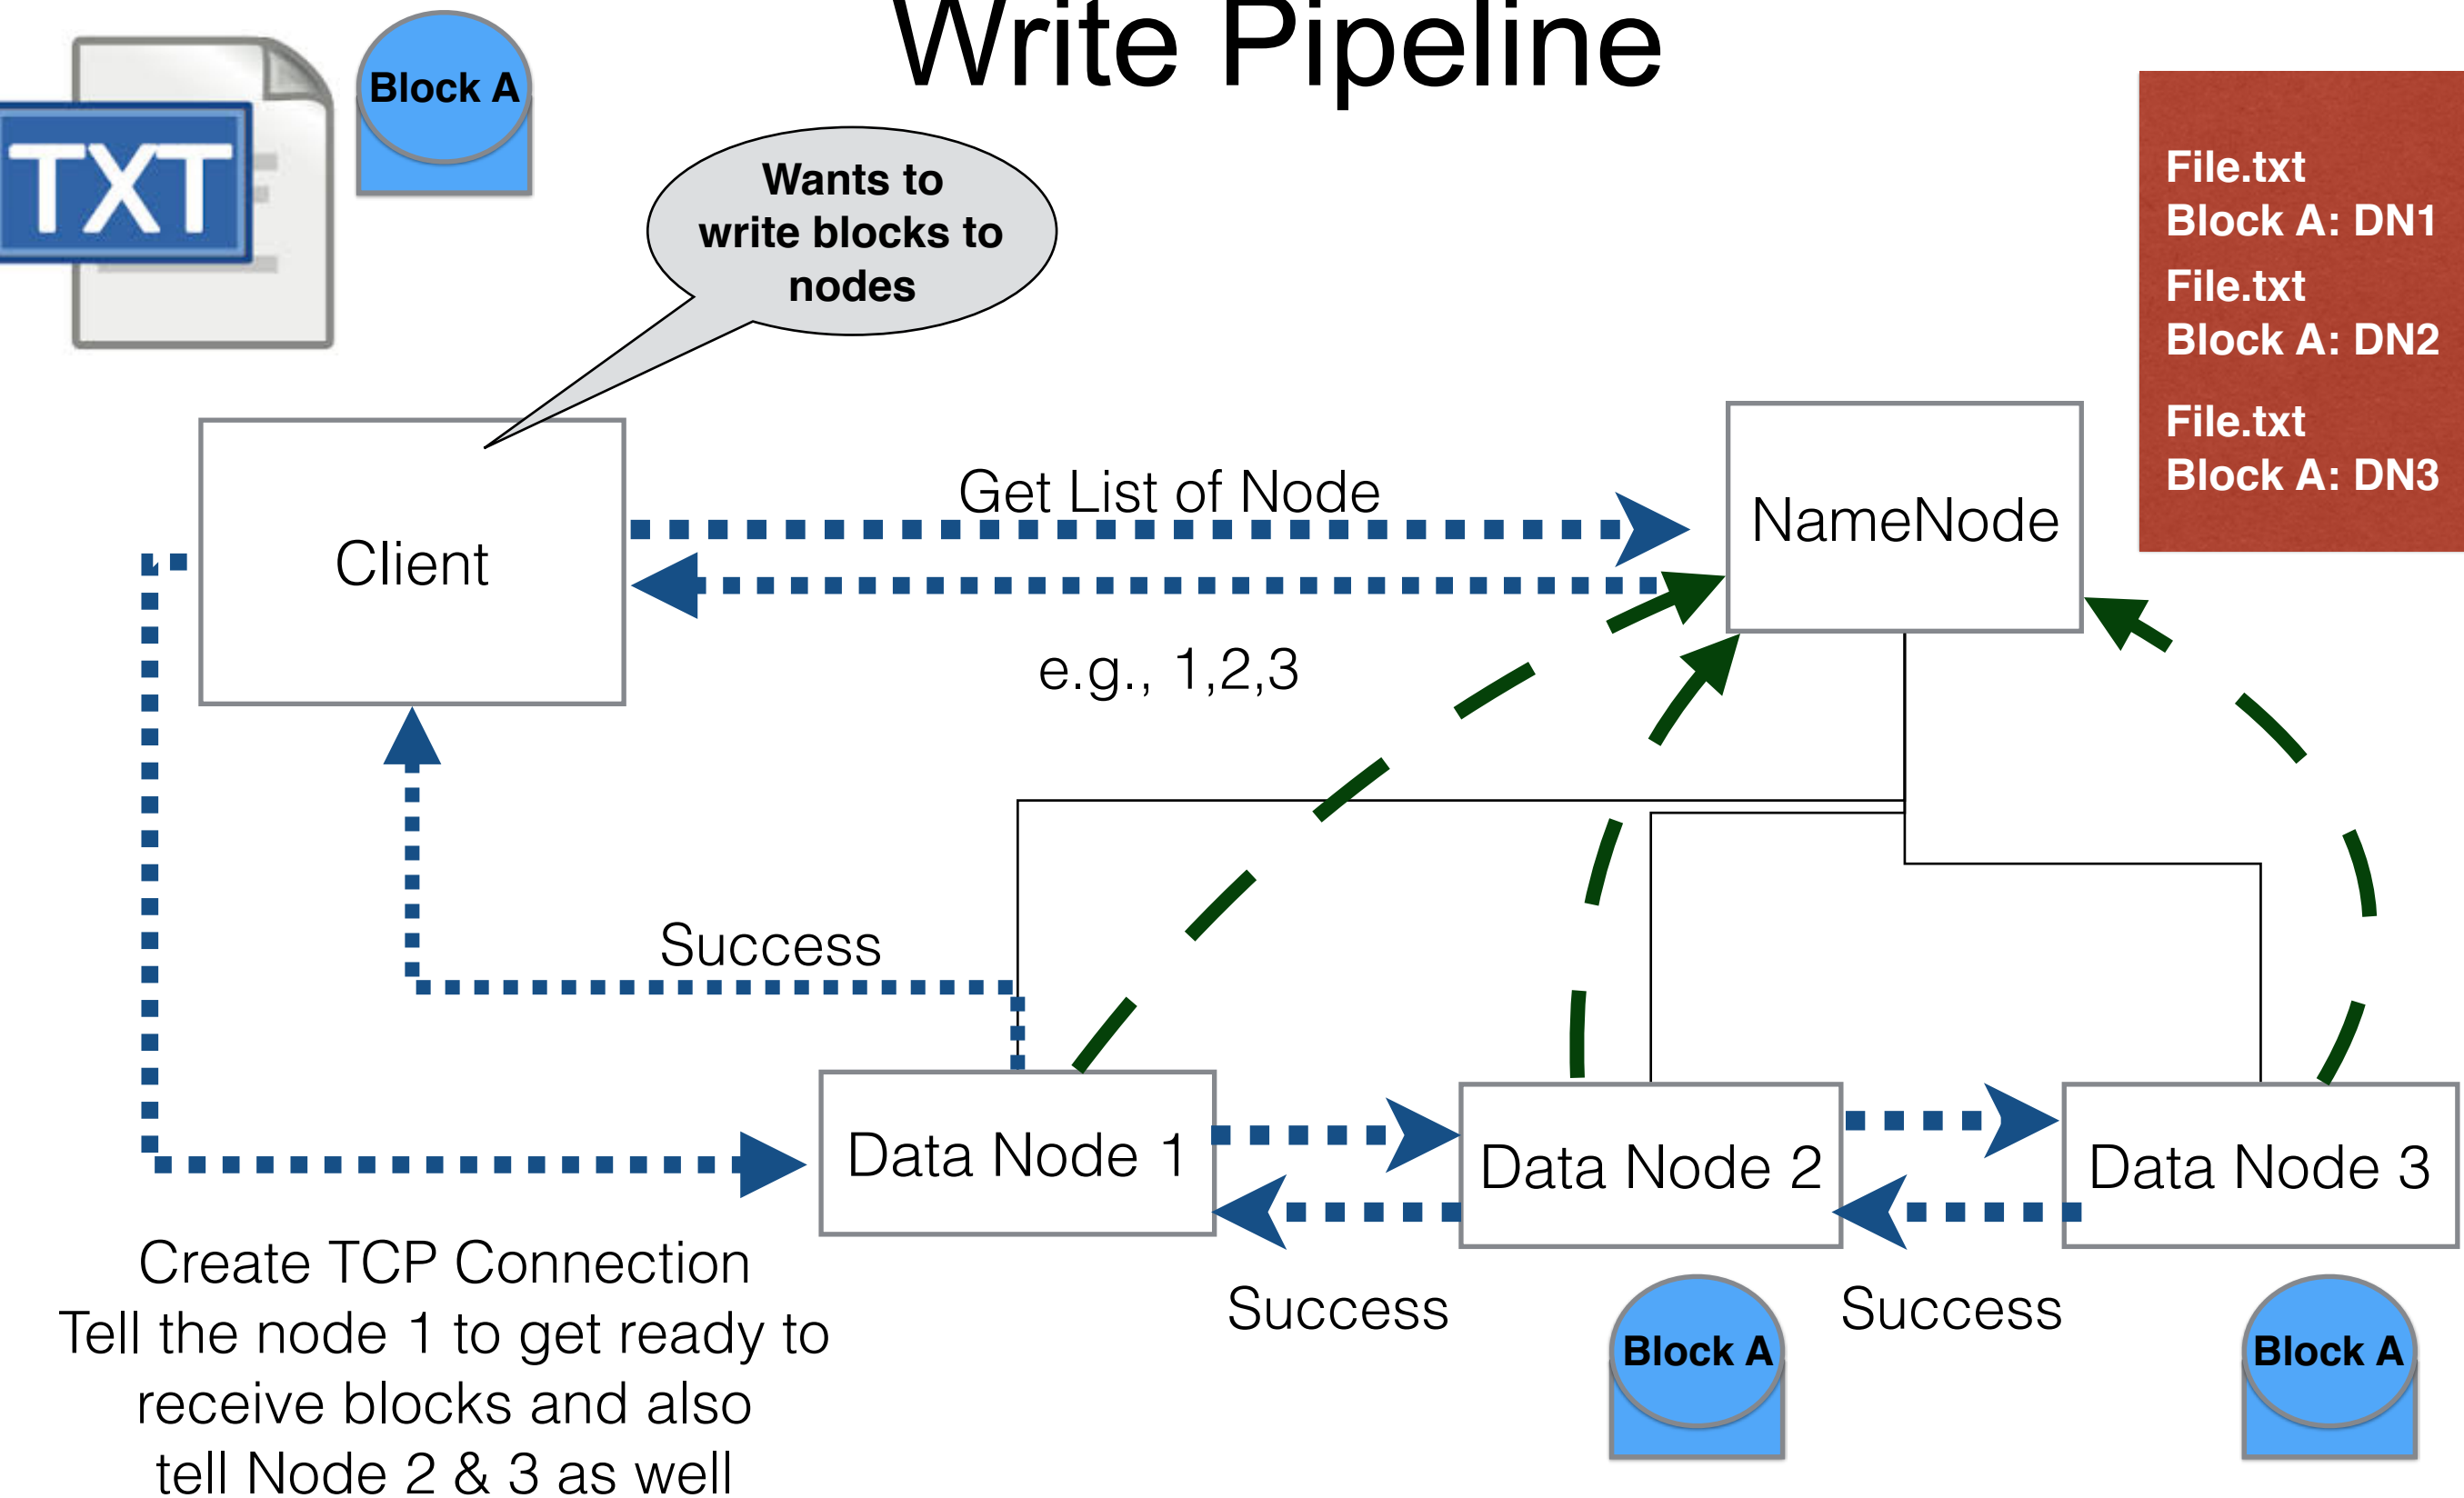
Data Storage

- Files are divided into blocks — Stored in multiple DataNodes
 - Blocks are configurable parameter in HDFS — default block size 128MB
- Replication Factor
 - Number of DataNodes a block is replicated
 - Default value is 3

Replication Factor —> 2




Write Pipeline



Rack Awareness

- In practical deployment nodes are deployed in racks
- Each block of data is replicated to multiple nodes to eliminate single point of failure
- Hadoop is rack aware:
 - Data loss prevention
 - What happens if all of the blocks are in the same racks?
 - What happens if racks goes down?



Read Pipeline

- Similar to the write pipeline
- Clients gets a list of data nodes for each block
- Establishes TCP connection with the first node
- Reads all blocks sequentially



Re-replication

- If the Name Node stops receiving heartbeats from a Data Node
 - Presumes it to be dead and any data it had to be gone as well.
- Based on the block reports it had been receiving from the dead node
 - Name Node knows which copies of blocks died along with the node
 - Make the decision to re-replicate those blocks to other Data Nodes.
 - Also consult the Rack Awareness data in order to maintain at least one copy in another rack
- If entire rack of servers falls off the network,
 - The NameNode tells remaining nodes in the cluster to re-replicate all of the data blocks lost in that rack.
 - Large amount of data could potentially traverse the network



Re-replication

- HDFS scan blocks for corruption.
- Every DataNode checks the block present in it
 - Verifies with the stored checksum, which is generated during the block creation.
 - In case block corruption is identified, NameNode is informed
 - NameNode marks the block in the DataNode as corrupt and initiates a re-replication of the block.



HDFS Federation

- NameNode has a single namespace and tightly coupled with DataNodes
 - All requests have to coordinate with NameNode to get blocks' location
 - Can easily become a bottleneck and limiting factor
- HDFS Federation enables multiple namespaces
 - Responsibility of NameNode is shared across multiple namespaces
 - Load balancer is used to route requests
 - Federated NameNode works in multi-tenant environment



HDFS Commands

- Very similar to unix
 - `hadoop fs <args>`
- `hadoop fs -mkdir <paths>`
- Example:
 - `hadoop fs -mkdir /usr/srafiqi/CloudComputing`
 - `hadoop fs -ls /usr/srafiqi/CloudComputing`
 - Put/Get:
 - `hadoop fs -put <localsrc> <HDFS destination path>`



HBase Overview

HBase

- Column-oriented NoSQL, Distributed Database designed for queries of massive data sets
- HBase can run in distributed and pseudo-distributed modes
- Data is stored in tables
- In pseudo-distributed mode
 - All HBase daemons runs on a single node.
 - Can run on local file system or an instance of HDFS
- HBase support auto-swarding
 - Tables are dynamically split and distributed by the database when they become too large



HBase

- In HBase, each key has a structure
 - Row key, column family and timestamp
 - Actual key-value pair mapping in HBase
 - (row key, column family, column timestamp)



Flexible Data Model

- HBase is a wide-column data store based on BigTable Concept
- Basic unit of storage in HBase is a table
 - Consists of one or more column families
- Data is stored in rows
 - Collection of key/value pairs
- Each row is uniquely identified by a row key
 - Row keys are created when table data is added
 - Row keys are used to determine the sort order and for data sharding



Flexible Data Model

- Columns may be added to a table column family as required without predefining columns
- Only the table and column families are required to be defined in advance
- No two rows in a table are required to have the same columns
- All columns in the column family are stored closer
- HBase is strongly consistent (not eventual)
 - Data is always consistent — but potentially increase latency
- Does not have the notion of data types
 - All data is stored as an array of bytes



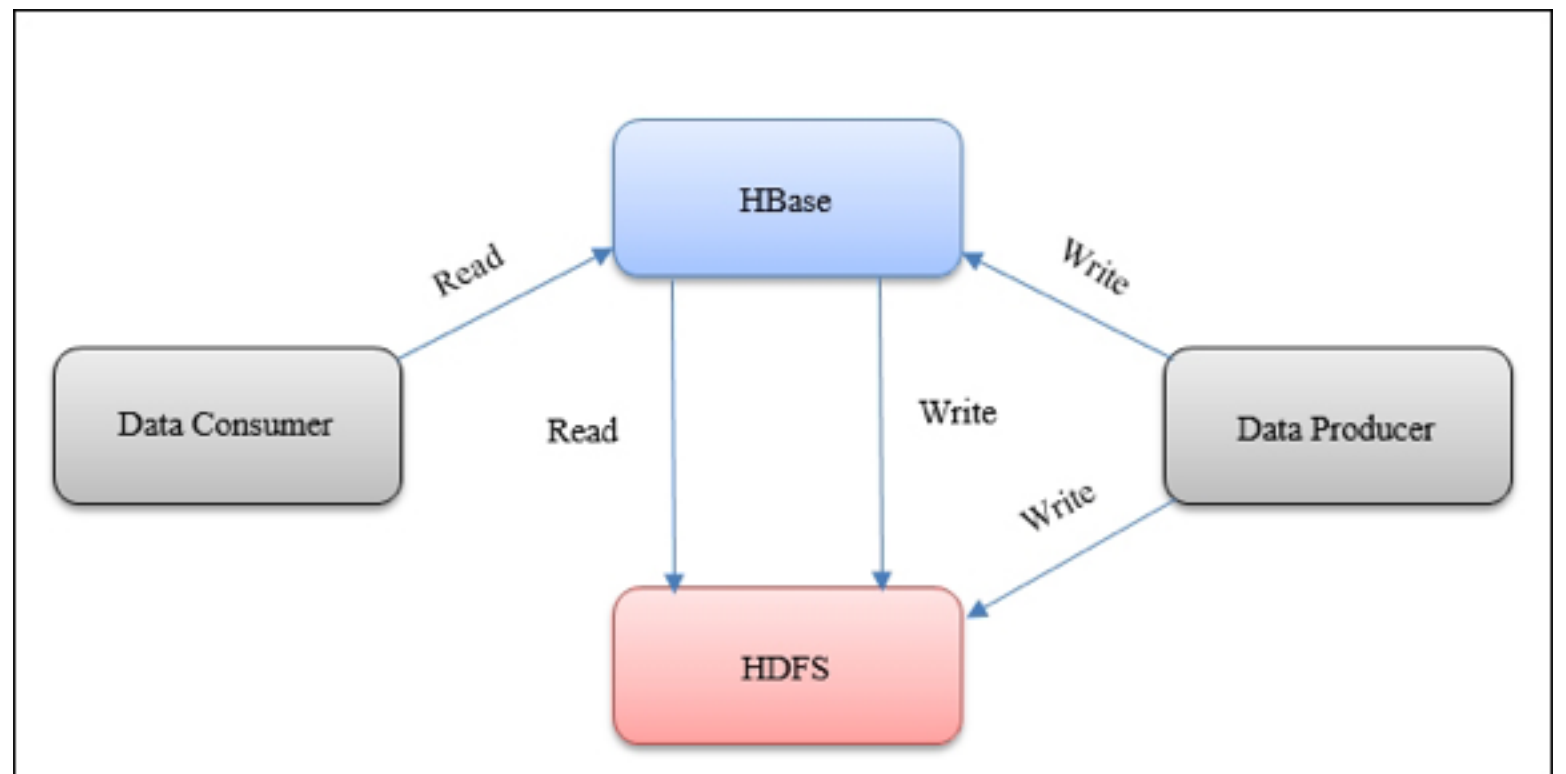
Scalable

- Basic unit of horizontal scalability is region
- Rows are shared by region
 - Region is a sorted set consisting of a range of adjacent rows stored together
- Table's data can be stored in one or more regions
- When a region becomes too large
 - Split into two at the middle row key into equal region



Transactional Capabilities

- HBase adds transactional capabilities to Hadoop
- Allows user to conduct CRUD operation in the DB
- Provides real-time random read/write access to data stored in HDFS



Main Features of HBase

- Linear and modular scalability
- Fault-tolerant storage for a huge amount of data
- Flexible data model
- Support for atomic operations
- Atomic sharding
- Easy-to-use Java API for client access
- Replication across the data center
- Strictly consistent read and writes
- Block cache and bloom filters for real-time queries



HDFS and HBase

| HDFS | HBase |
|--|--|
| HDFS is a distributed file system well suited for the storage of large files. It is not a general-purpose filesystem and does not provide fast individual record lookups in files. | HBase is built on top of HDFS and provides fast record lookups for large tables. It internally puts the data in indexed StoreFiles available in HDFS for fast lookups. |
| It is suitable for high latency operations batch processing. | It is built for low latency operations. |
| The data here is mainly accessed through MapReduce. | It provides access to single rows from billions of records. |
| It has been designed for batch processing, so it does not have the concept of random read and writes. | It supports random read and writes. |



Data Model

- **Table** — All data is stored there — Logical collection of rows
- **Rows** — Data is stored in rows, rowKey (unique identifier) but has no data type
- **Column Family** — Columns are grouped in column families.
 - All columns of column family has same prefix (c1)
 - All the columns in column family are stored in the File. Required to declare column families upfront
- **Column Qualifier** — In each column family, data is addressed in a column identifier
- **Timestamp** — Version in HBase. Whenever a value is written, it is stored with a version number
 - Version number is a timestamp of when cell was written

| Table X | | | | |
|---------|---------------|------------------|--------------|---------------------|
| rowkey | Column Family | Column Qualifier | Timestamp | Value |
| x | C1 | “foo” | “1223455634” | 5 |
| | | | “1245564567” | “Hadoop” |
| | | “test” | “1235674323” | 34 |
| | | | “1256348945” | 564 |
| | C2 | “1.0002” | “3467893423” | “Guide” |
| | | | “1235576347” | “Yet another value” |
| y | C2 | “2-12-2014” | “1235752574” | “Completed Task” |
| | | “CValue” | “1235573212” | 3576 |
| | | “qualifier” | “1245322456” | “My Data” |

Accessing data

- Four primary data operations that we can execute to access data in HBase:
 - Get, Put, Scan, and Delete
- Get — Returns an attribute for a specified row.
- Put — Used to either add a new row in the table or update an existing row.
- Scan — Allows to iterate over multiple rows for a defined attribute.
 - Using scan, user can match columns with filters and other attributes, such as the start and end time.
- Delete — This is used to remove a row from a table.

