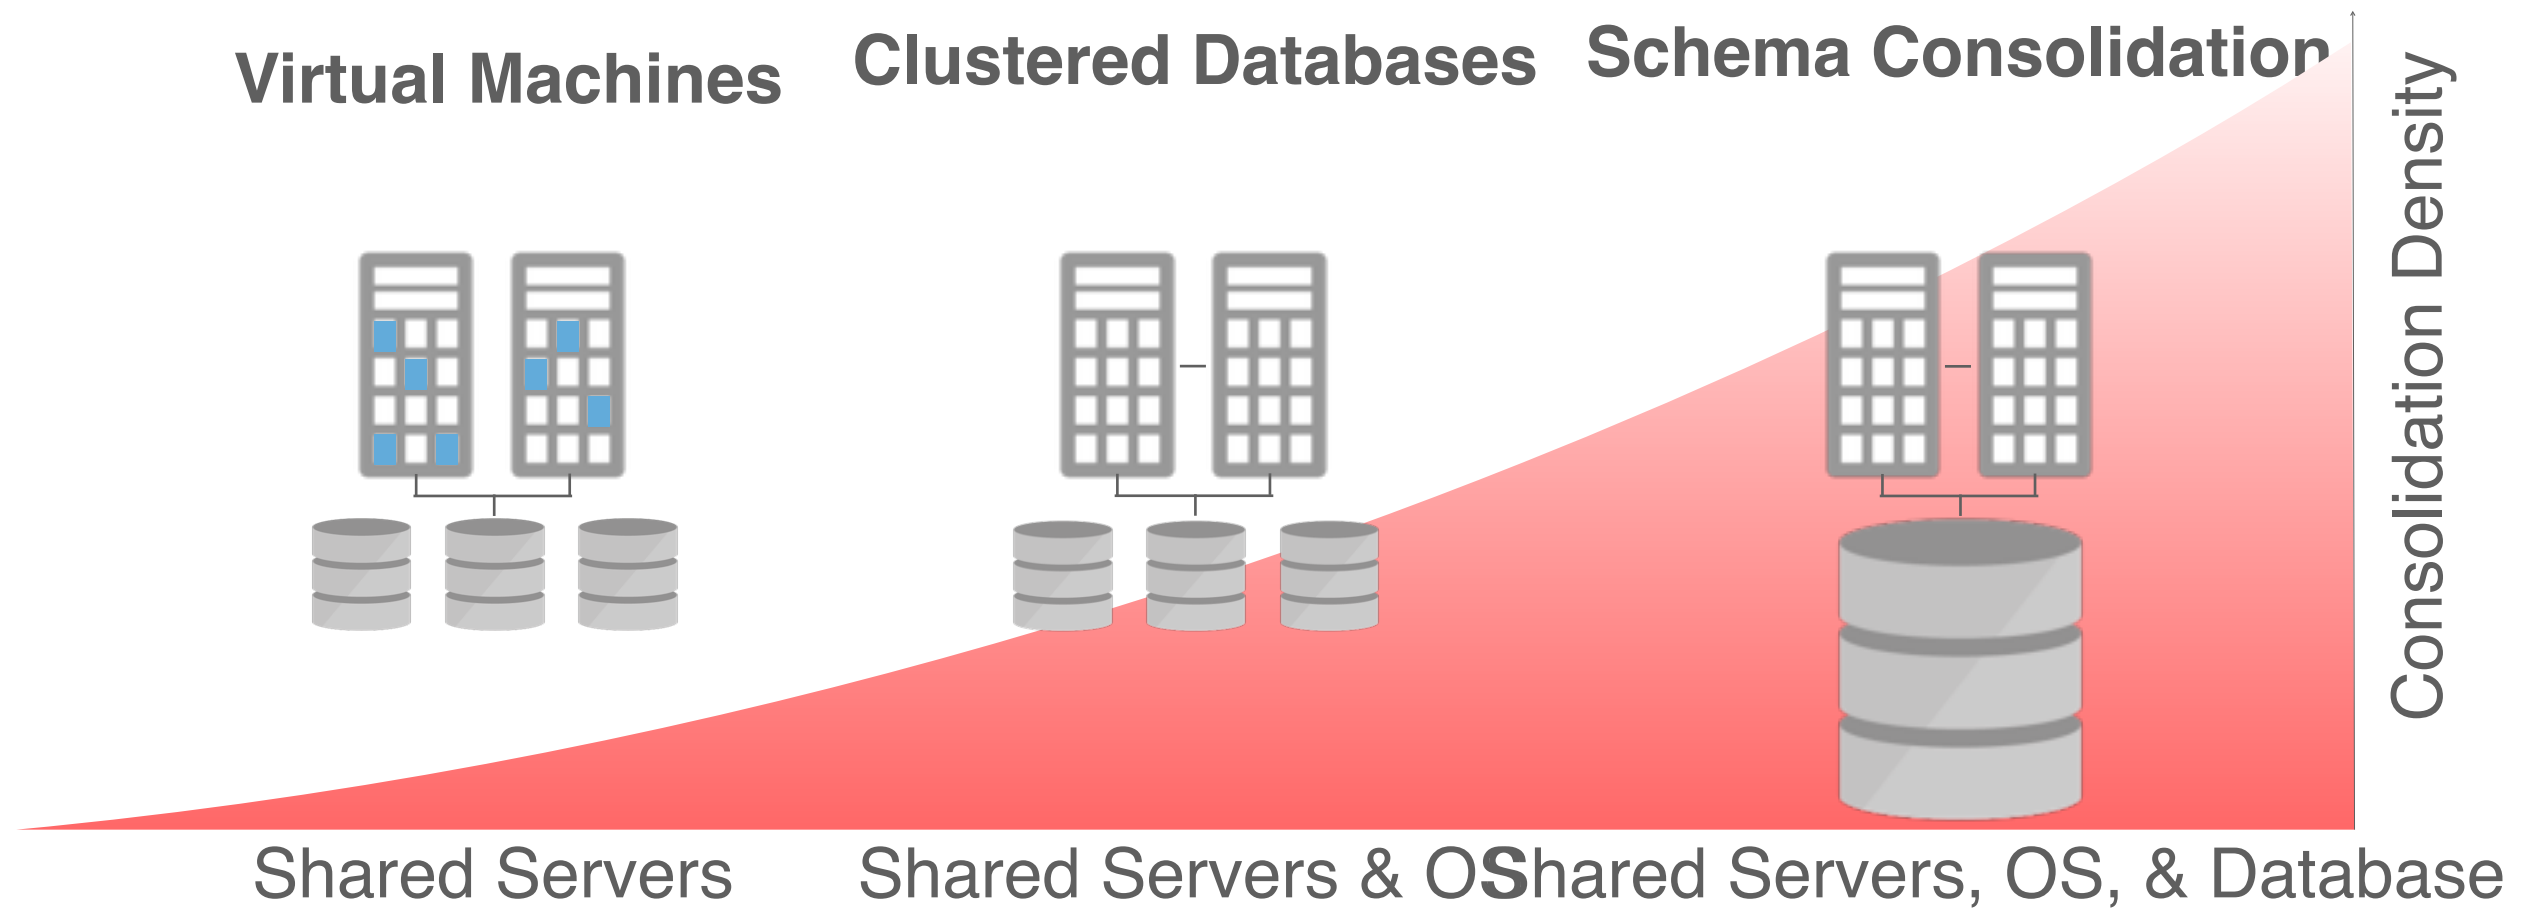



Multitenant Databases

Traditional Consolidation Methods



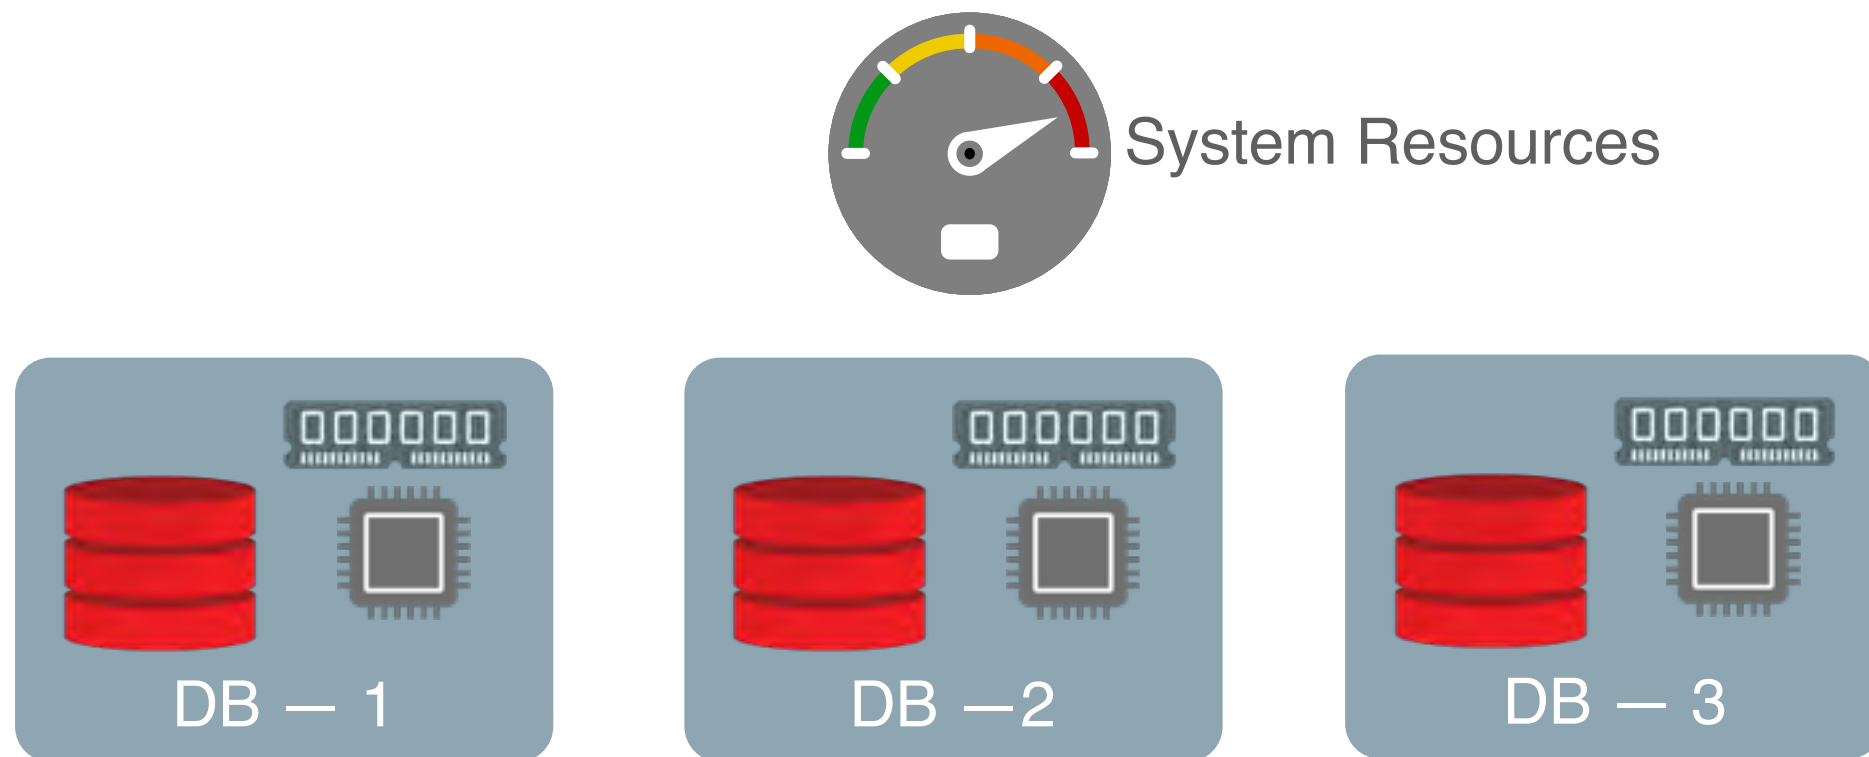
Consolidation through Multitenancy

Simplifies consolidation; enables Database as a Service




Traditional Database

Requires memory, processes and database files







Multitenant

Maximize the resource utilization



System Resources




Architecture

Components of a Multitenant Container Database (CDB)



Pluggable Databases




Multitenant Container Database



Pluggable Datababasese is a Portable Data

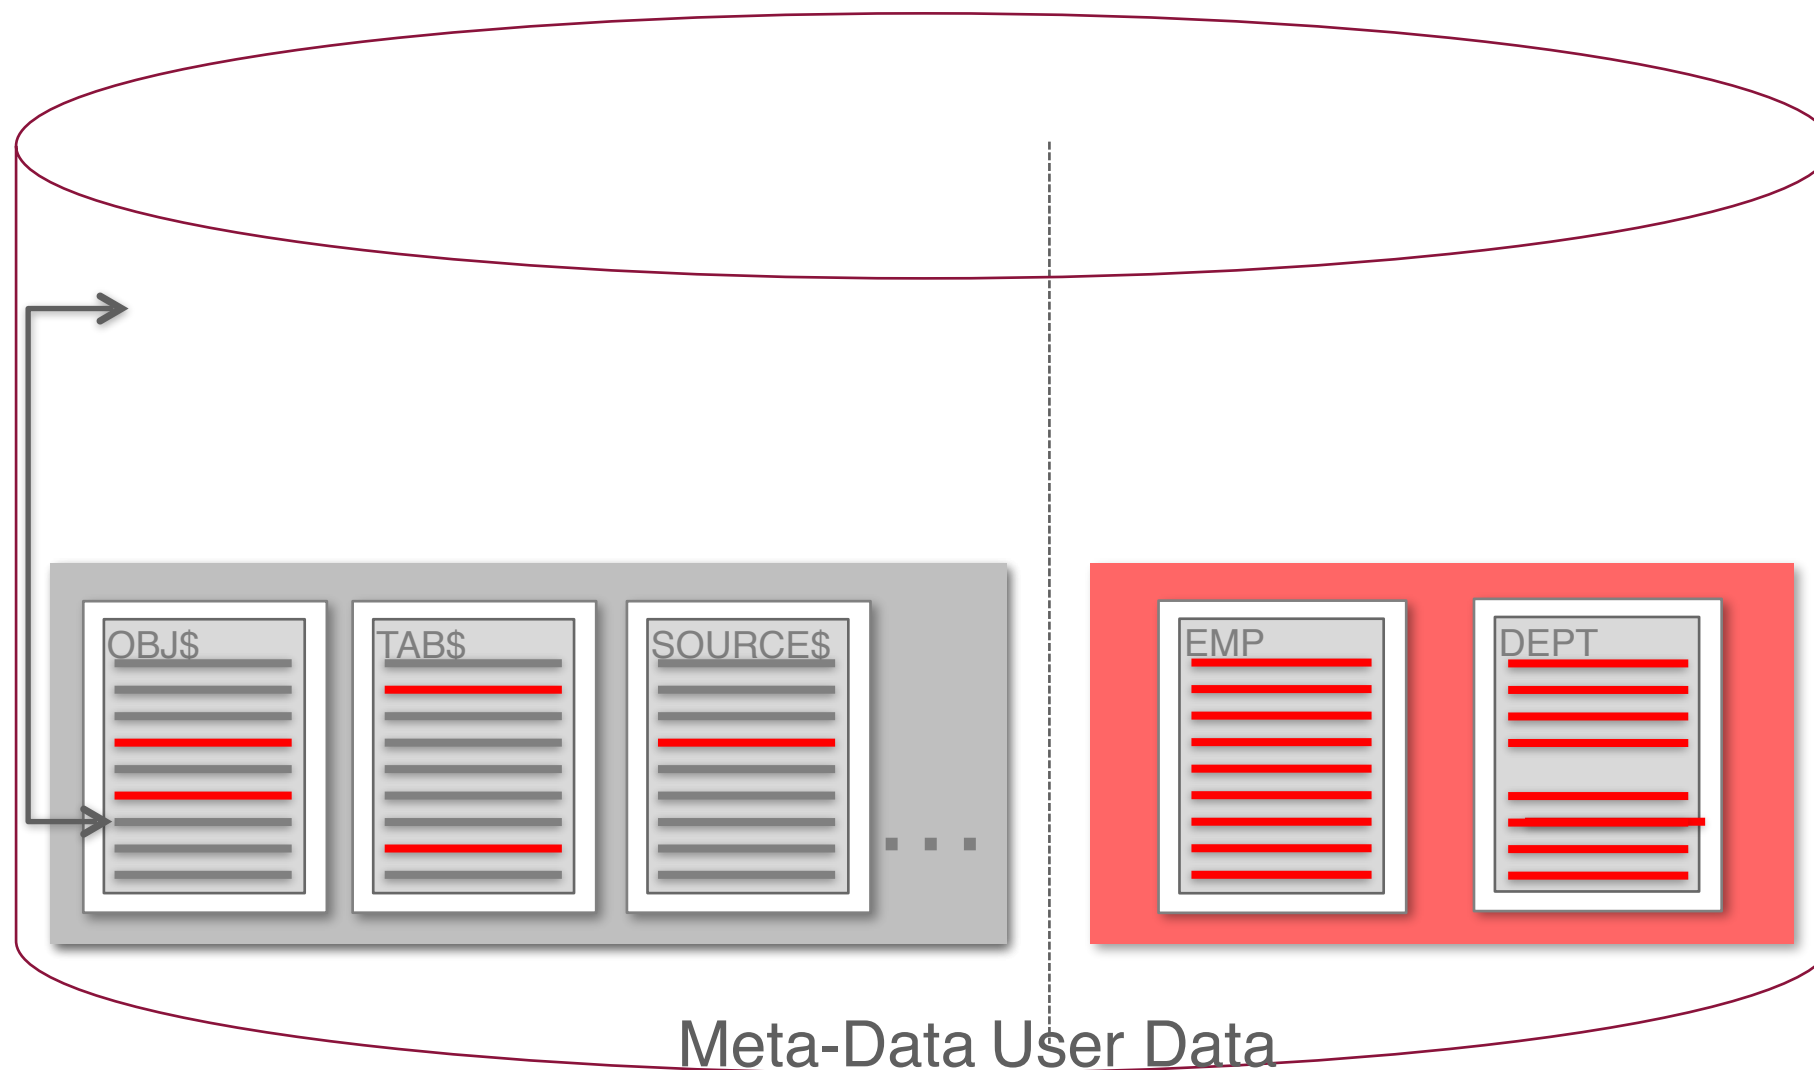
Unplugging a PDB and plugging it in



- Simply unplug from the old CDB...
- ...and plug it into the new CDB
- Moving between CDBs is a simple case of moving a PDB's metadata
- An unplugged PDB carries with it lineage, opatch, encryption key info etc.



Data and User Data



- New database contains Oracle meta-data only
- Populate database with user data
 - Oracle and customer meta-data intermingled
 - Portability challenge!
- Multitenant fix:
Horizontally-partitioned data dictionary
 - Only Oracle-supplied meta-data remains in root



Unplug/Plug using SQL

Unplug

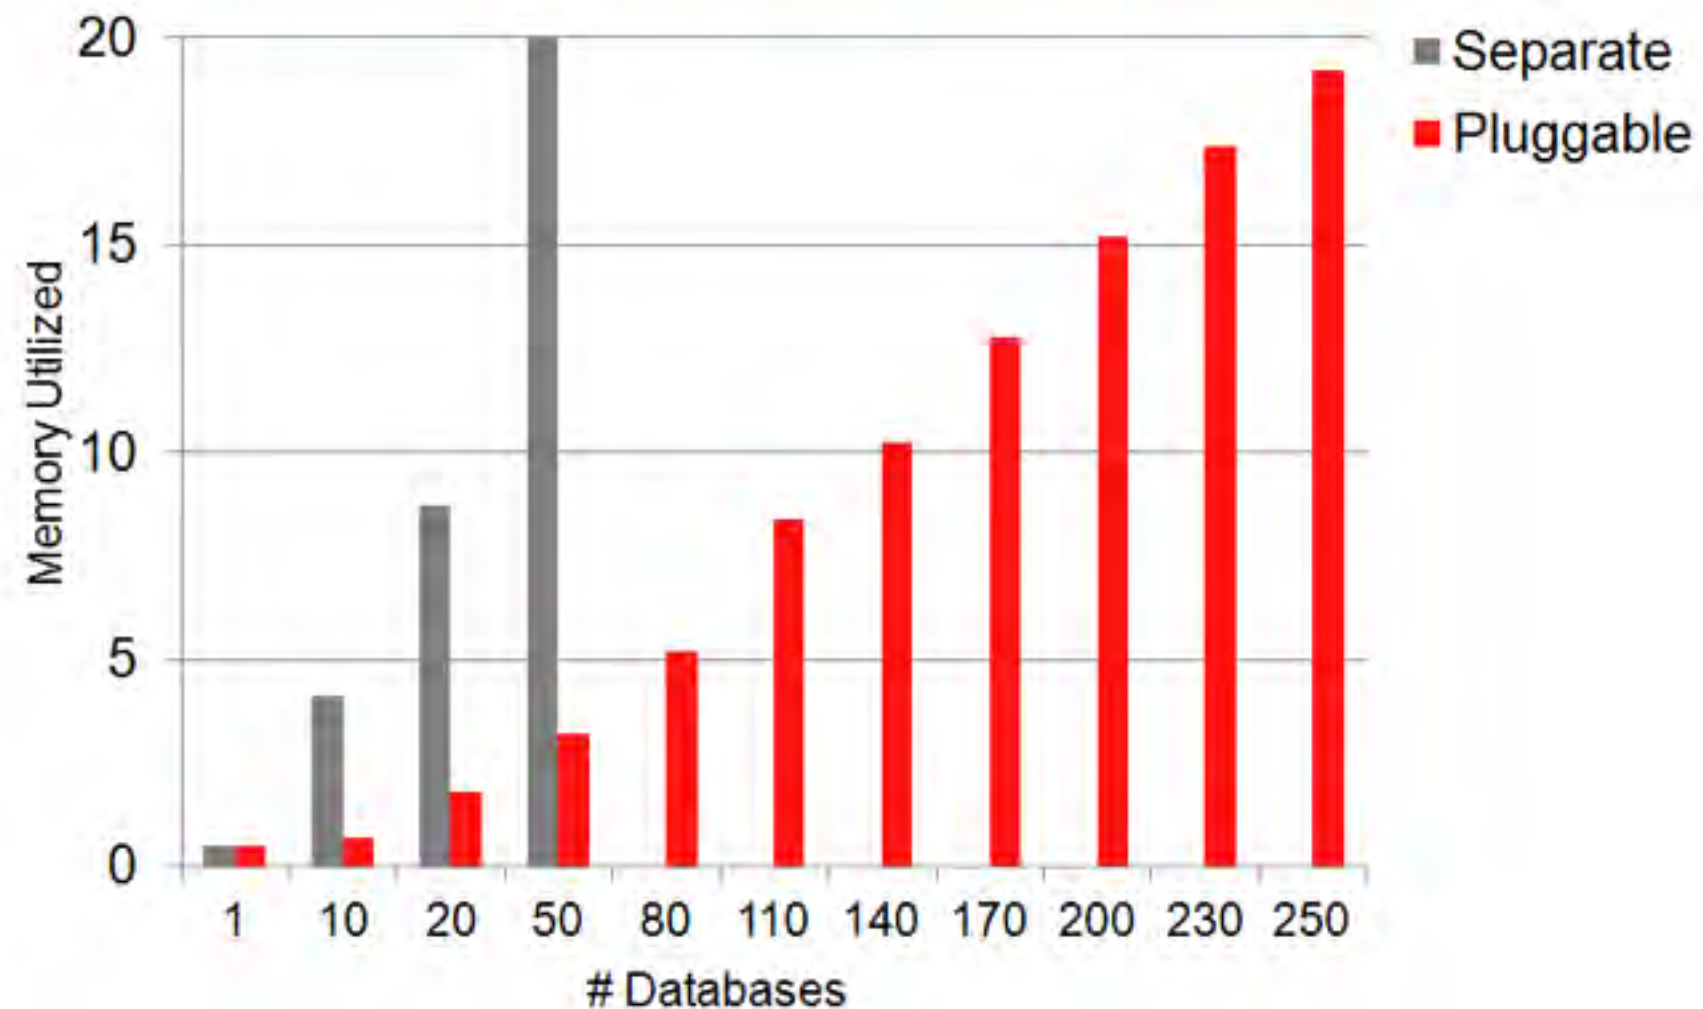
```
alter pluggable database OE  
unplug into '/u01/app/myDB/data1/.../  
oe.xml'
```

```
create pluggable database My_PDB  
using '/u01/app/data2/.../oe.xml'
```



Multitenant vs Separate Database

- OLTP benchmark comparison
- Only 3GB of memory vs. 20GB memory used for 50 databases
- Multitenant architecture scalable to large # of databases



Files in Container

- Each PDB has its own set of tablespaces including SYSTEM and SYSAUX
- PDBs share UNDO, REDO and control files, (s)pfile
- By default the CDB has a single TEMP tablespace but PDBs may create their own



Users

- Local users are the successors for customer-created users in a non-CDB
- A local user is defined only in a PDB
- A local user can administer a PDB
- A common user is defined in the root and is represented in every PDB
- A common user can log into any PDB where it has “Create Session” and can therefore administer a PDB
- The system is owned by common users



Common Users and Privileges

- A common user can be granted privileges locally in a PDB (or root) and therefore differently in each container
- A common user can, alternatively, be granted a system privilege commonly – the grant is made in root and every PDB, present and future
- You can create a common role
- A common role can be granted to a common user commonly
- Authorization is checked in the container where the SQL is attempted considering only the privileges that the user has in that container



Granting the SysDBA privileges in a PDB

- When a user authorizes AS SYSDBA in a PDB, the effect is contained to within that PDB
- This does not enable the user to get to another container
- Getting to a new container depends only on the privileges that the user has in the new container

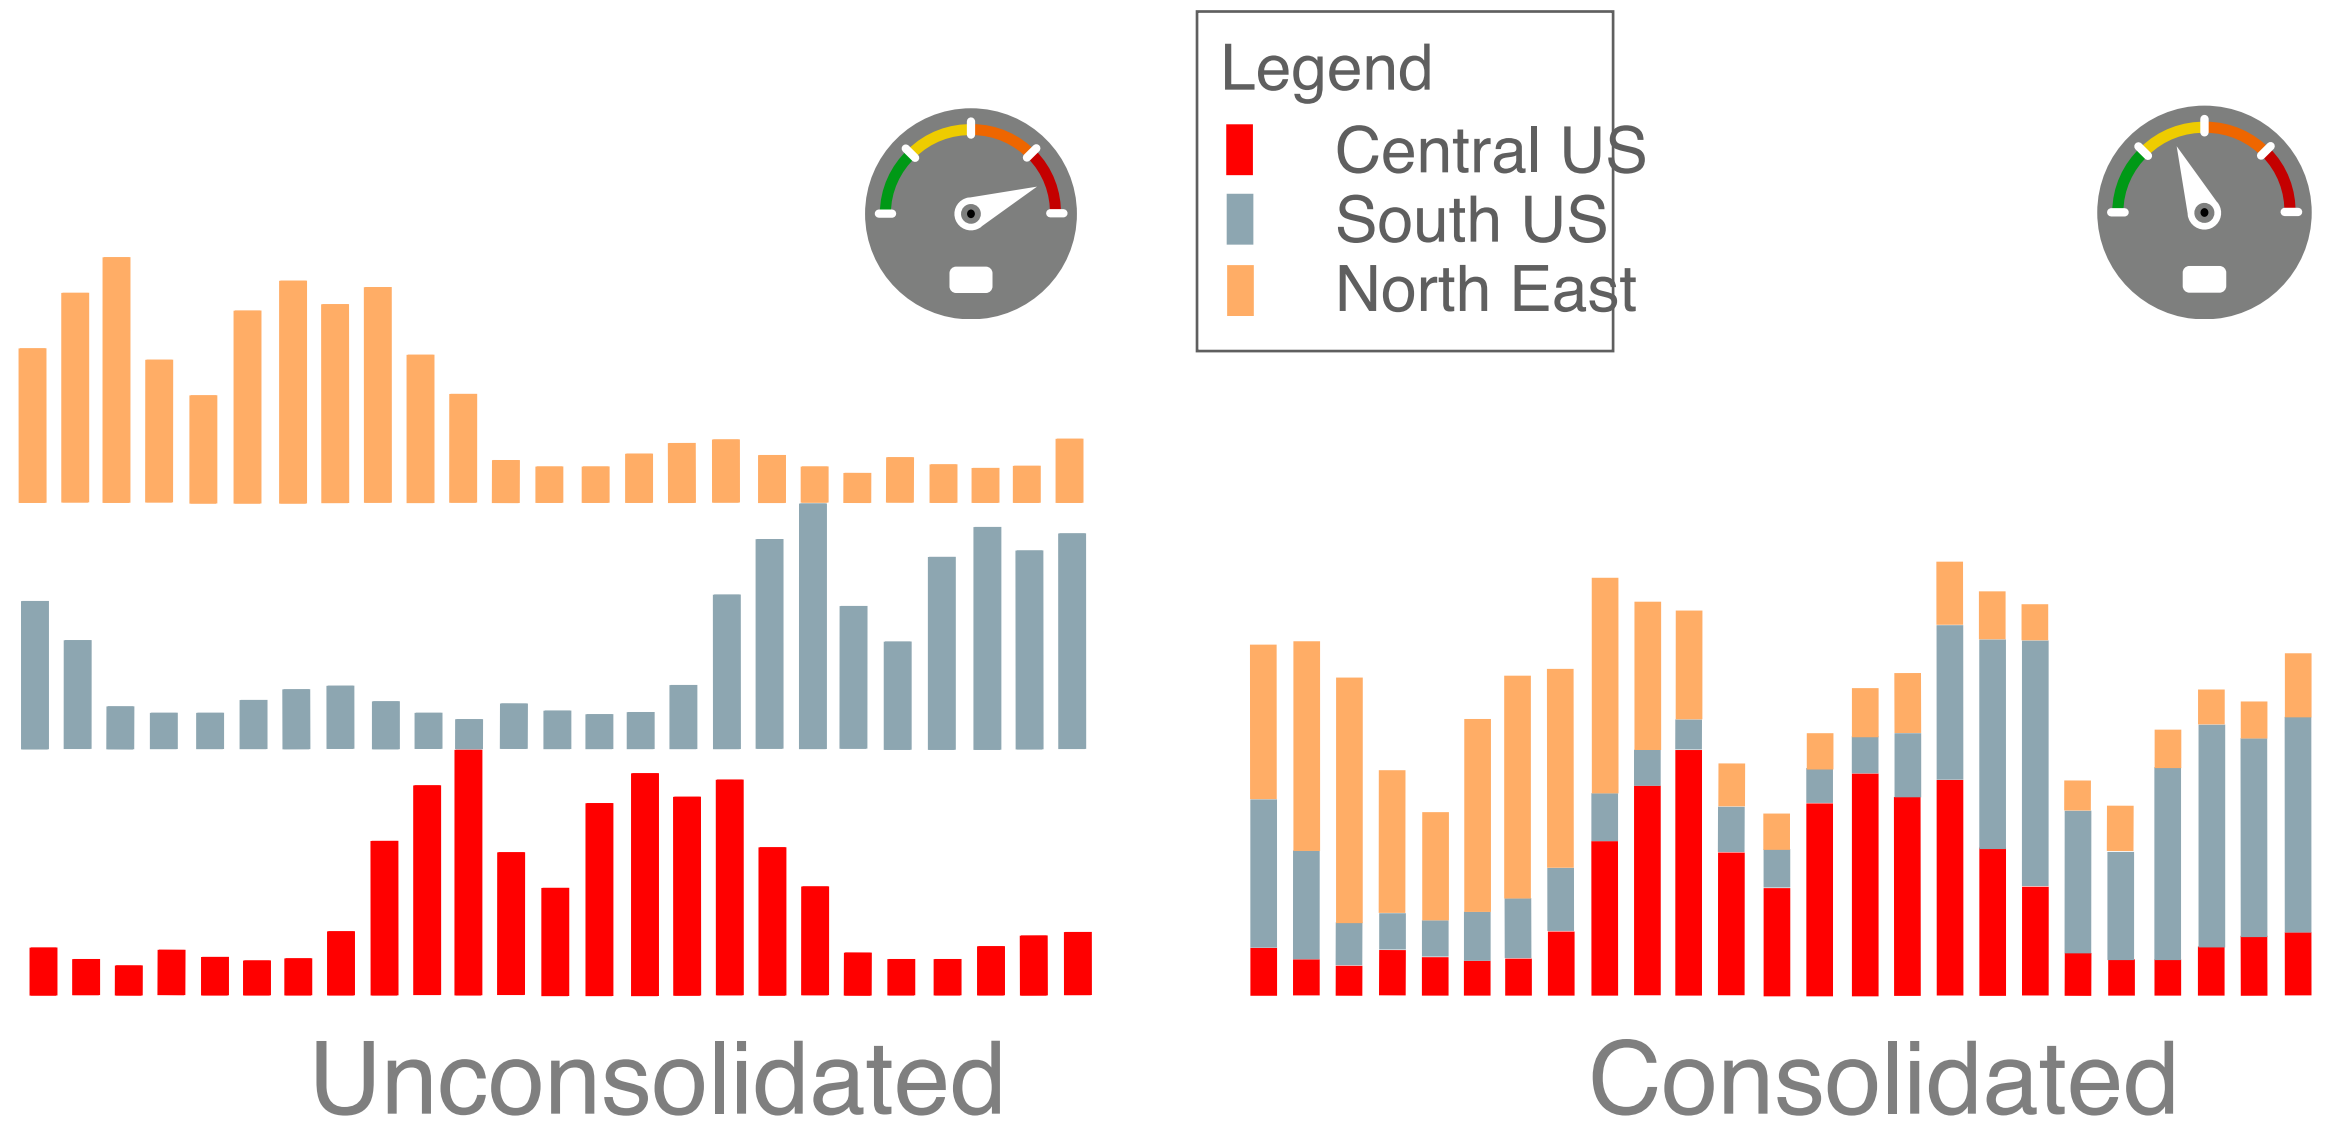


Creating Common Users

```
create user c##Über_Administrator  
identified by pwd container=all
```



Benefits of Consolidation



Consolidation Math

- Simple Example:
- 5-core server
- 60% steady load
 - 2 cores “unused”
- Background overhead 1 core
 - 2 cores required for application processing




Consolidation Math Cont.

- Consolidation Server
- 10 core
- Target utilization 70%
 - 3 cores unused
- Overhad 1 core per CDB
 - 6 core for app processing
 - 2 core required for each application (previous step)
- Multitenant Consolidation
 - 2x the core count
 - 3x the application



Manage Many Databases as One

Backup databases as one; recover at pluggable database Level



Manage Many Databases as One



Simplified Patching and Upgrades

- Patches and upgrades are applied at multi tenant container level
- All PDBs in that containers are automatically upgraded
- Easy pathway to fallback as well



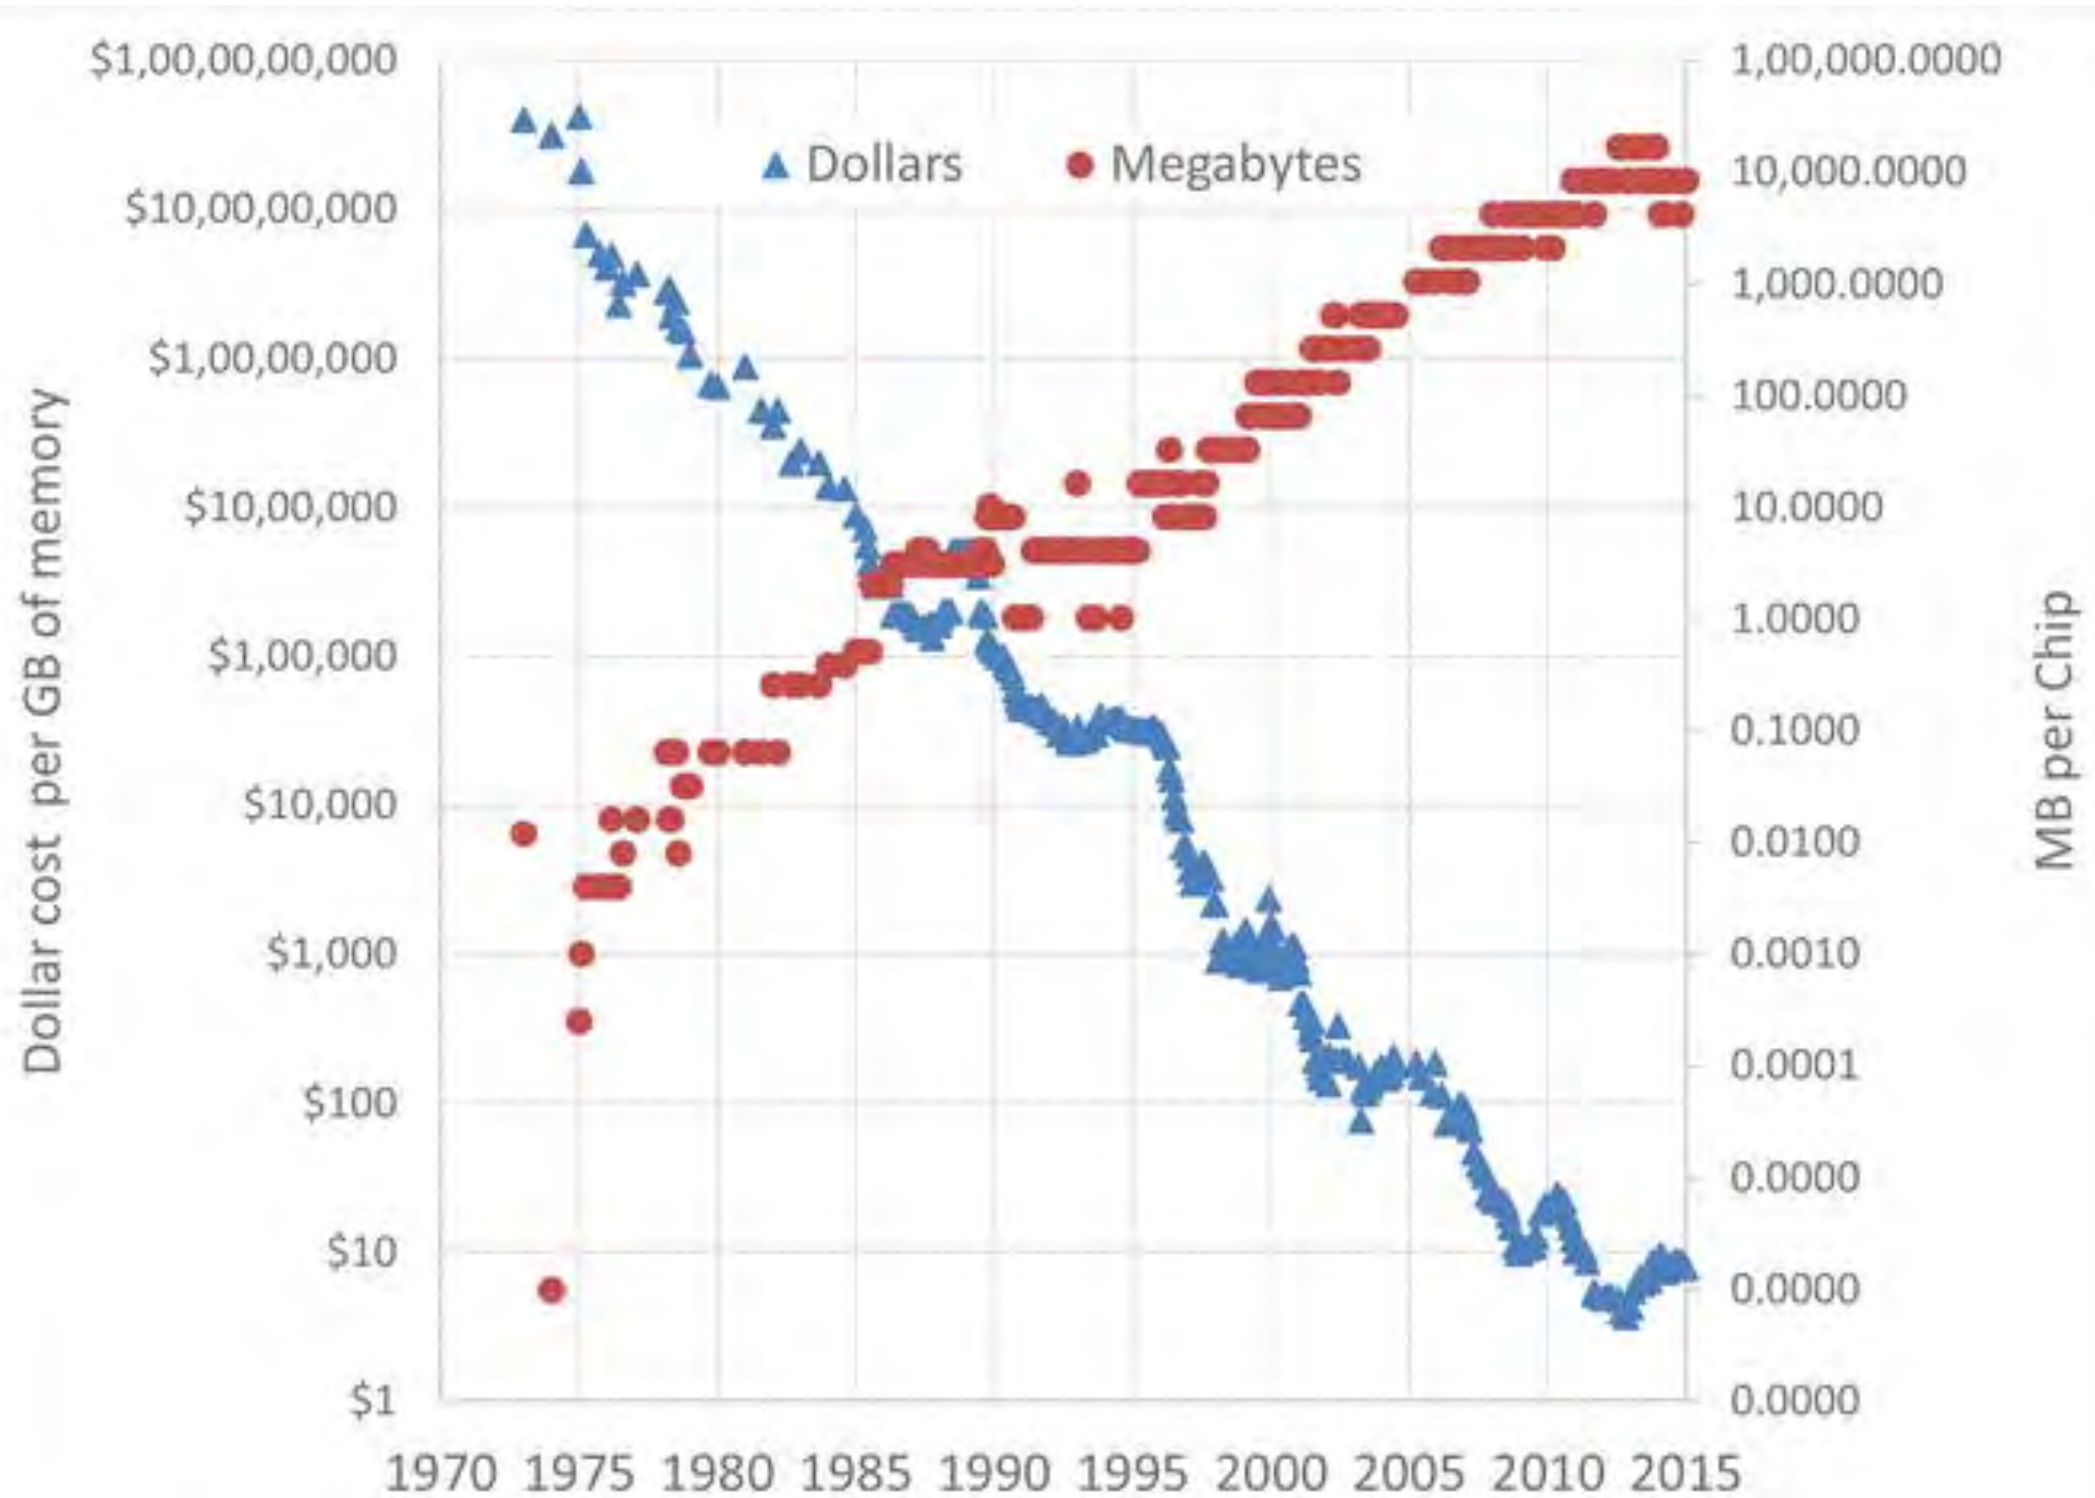
In-Memory Databases

Traditional Databases

- Traditional Database uses memory to cache data stored on disk
 - Generally show significant performance improvements as the amount of memory increases
 - Some database operations must write to a persistent media
 - COMMIT requires a write to a transaction log
 - Periodically database writes checkpoint blocks from memory to disk
 - Taking advantage of large memory requires an architecture that is aware of memory residence of data



In-Memory Databases



Source: <http://www.jcmit.com/memoryprice.htm>

Changes to Traditional Databases

- Cache-less architecture
 - There is no point caching in memory what is already stored in memory
- Alternative persistence model
 - Data in memory disappears when the power is turned off
 - Database must apply alternative mechanism for ensuring that no data loss occurs.



In-Memory Databases

- TimesTen
- SAP HANA
- Redis
- RDBMS with In-Memory options



TimesTen

- One of the earliest implementation of in-memory database system
 - RDBMS compatible
- All data is memory resident
 - Persistence is achieved by periodic snapshots of memory to disks and
 - Writing to disk-based transaction log following a transaction commit.

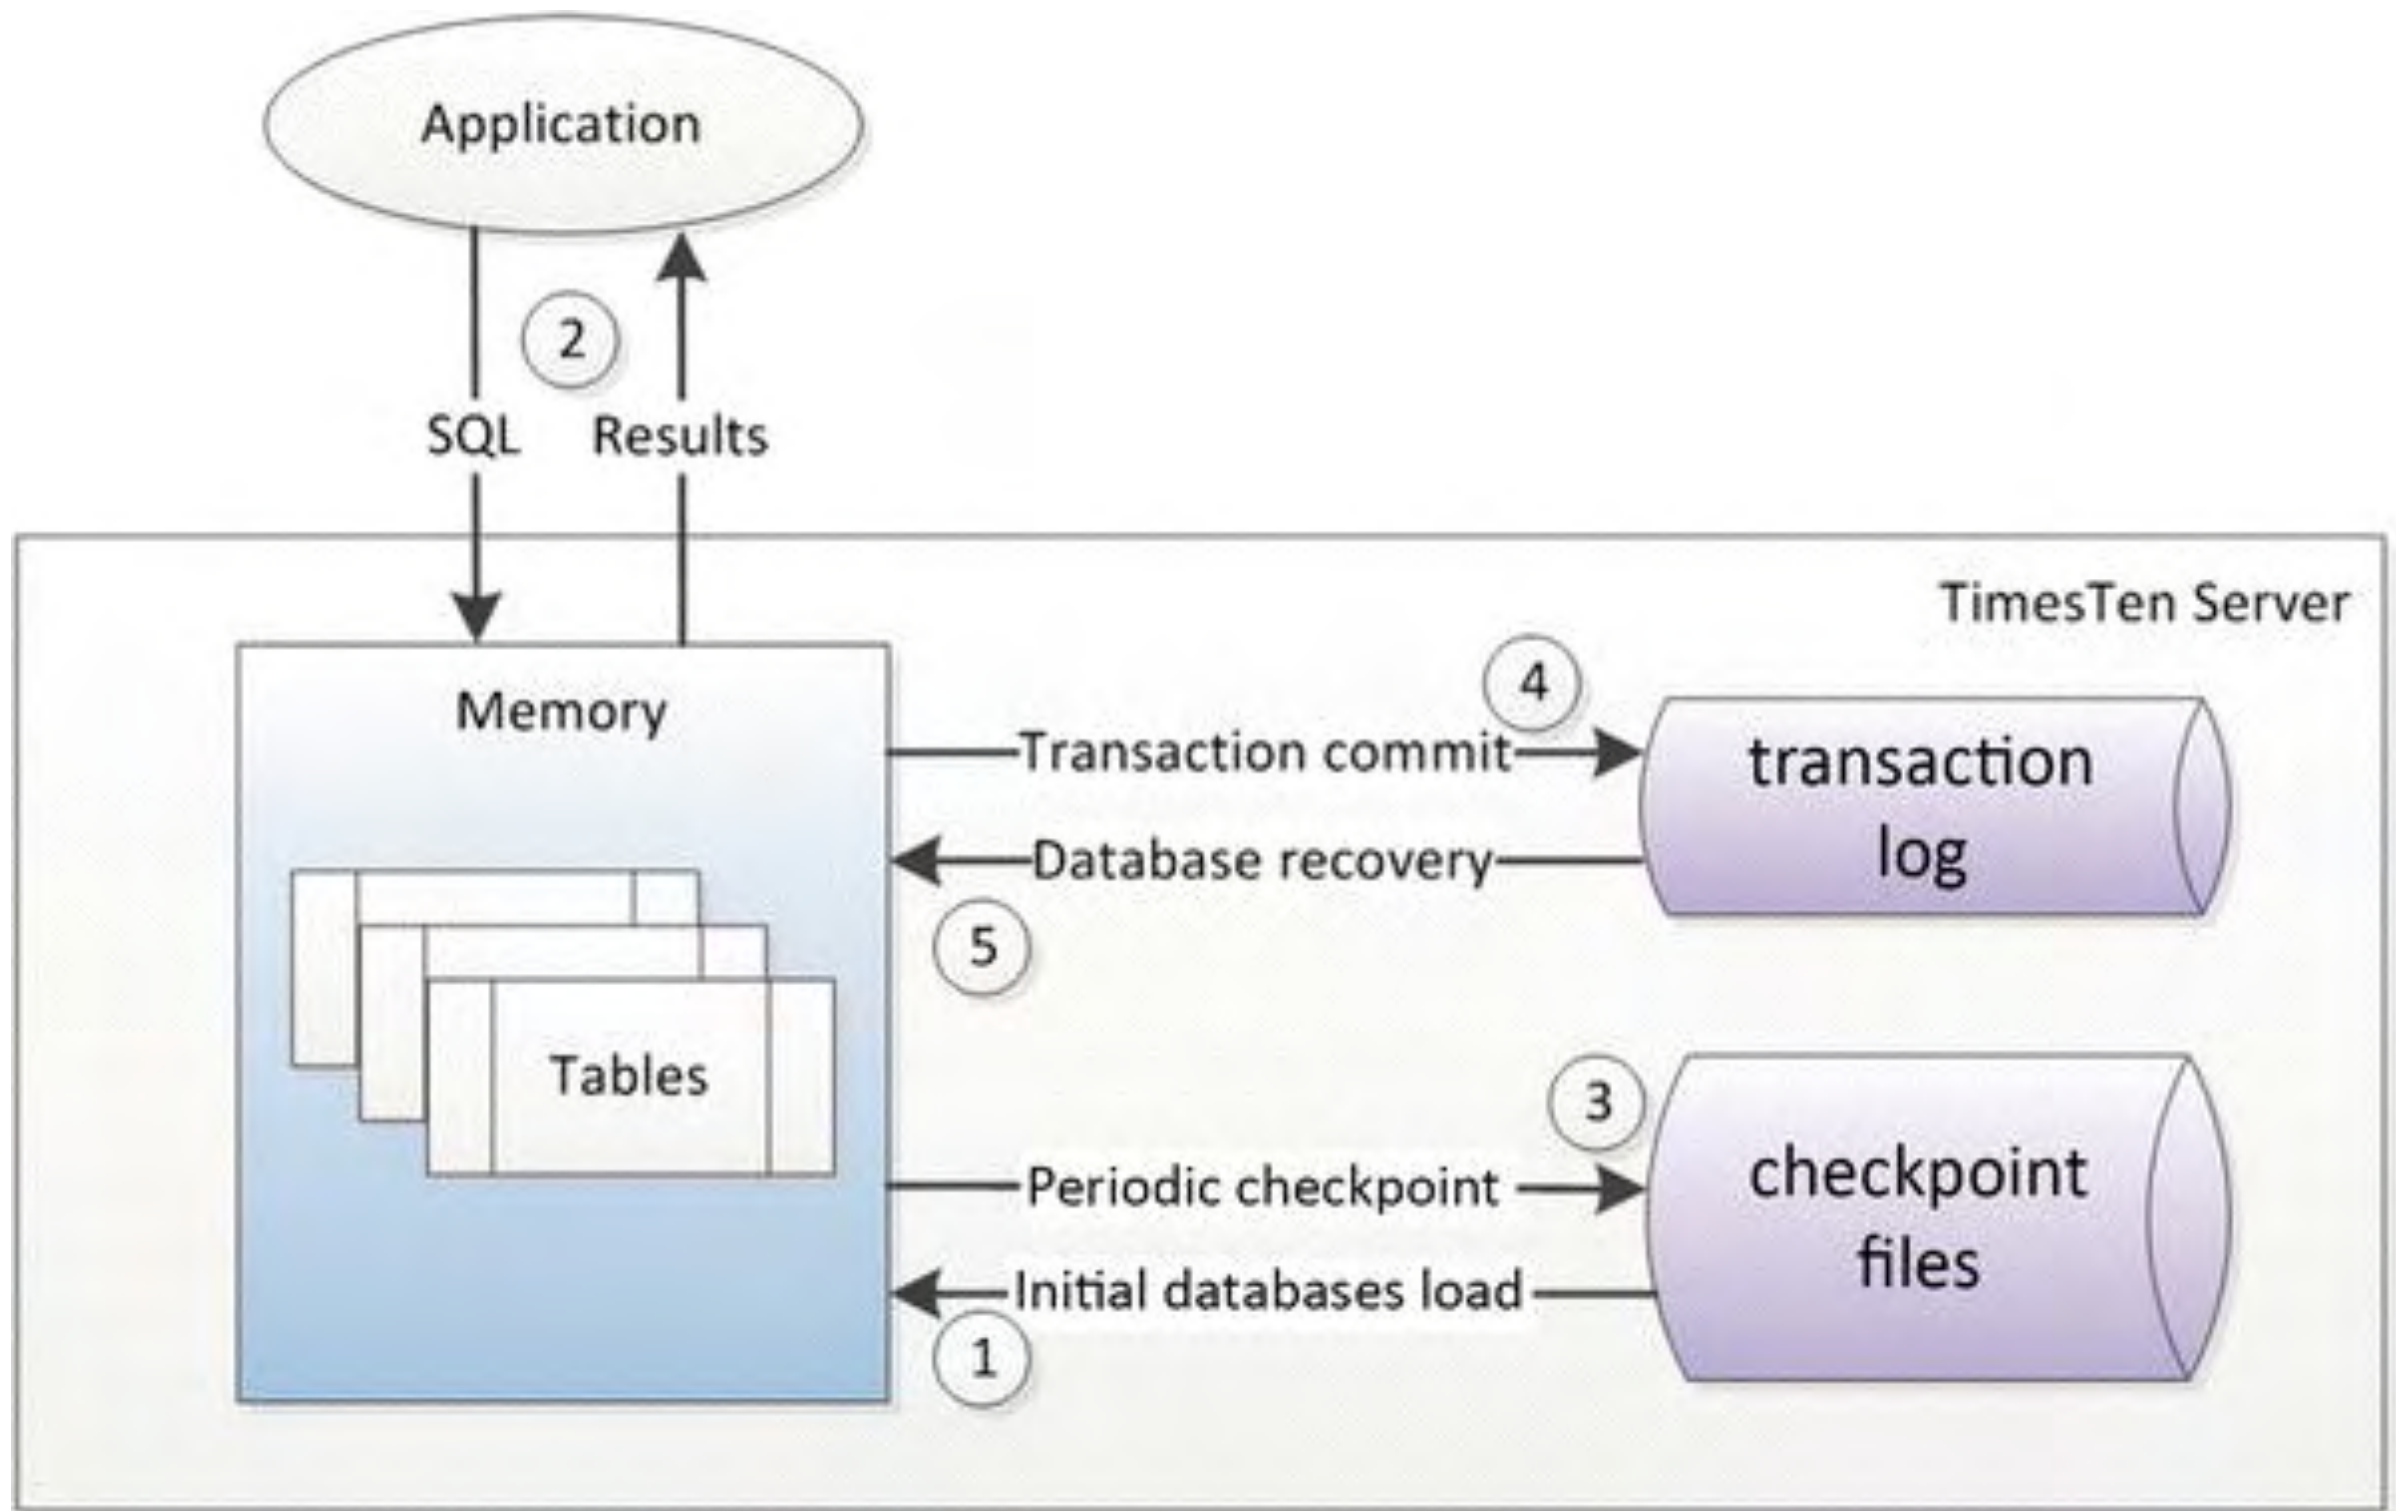


TimesTen Architecture

- In default configuration
 - Disk writes are asynchronous
 - If the power fails between transaction commit & transaction log
 - Data could be lost — Transaction durability is not guaranteed
- Synchronous writes — DB operation must wait on disk



TimesTen Architecture

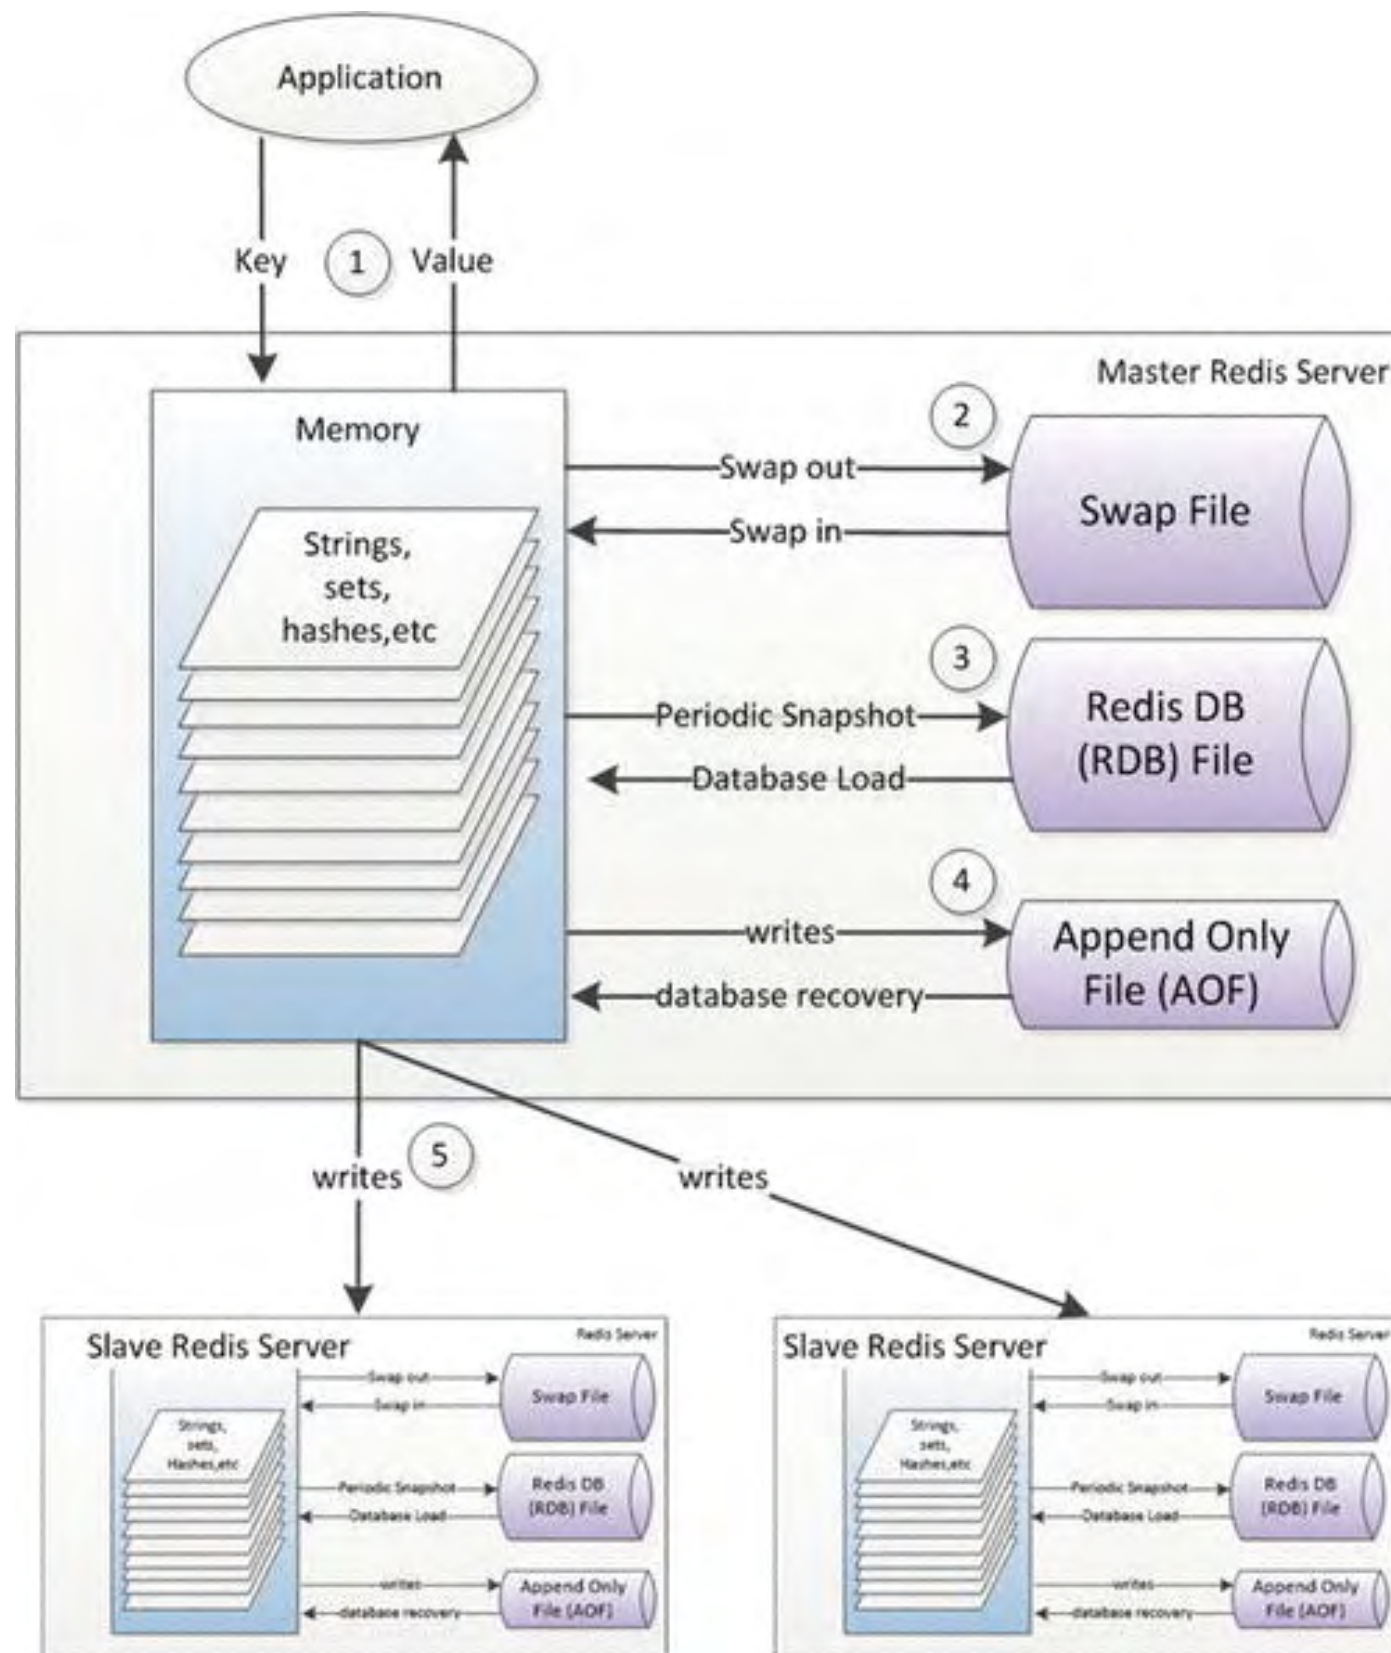


Redis

- In-Memory key-value pair store
- Possible to operate datasets larger than available memory
 - Virtual memory feature — Redis swap out older key value to a disk file
- Redis uses disk file for persistence:
 - Snapshot — stores copies of the entire Redis system at a point in time
 - Append Only File (AOF) — Keeps journal of changes that can be applied in case of a failure



Redis Architecture



SAP HANA

- In-Memory designed for Business Intelligence
- Tables can be configured row-oriented or column-oriented
 - Data suited for BI —configure column-oriented
 - OLTP — Row oriented.
- All row store is guaranteed to be in memory
- Column store is by default loaded on demand
 - Can be configured for immediate loading on startup
- Persistence architecture of HANA uses snapshots and journal file patterns similar to TimeTen and Redis
- ACID is enabled by redo log
 - Redo log is written upon transaction commit
 - To speed up transaction — redo logs are placed on solid state disk



VoltDB

- Designed not to wait for disk IO
- Designed with explicit intention to not requiring disk IO
- Support aCID transactional mode through replication across multiple machines
 - Transaction commit only complete once data is successfully written to more than one physical machine




Column & Row-store

- OLTP
 - Row format —
- Analytic queries
 - Columnar format



In-Memory Columnar Option




***ALTER TABLE sales
INMEMORY;***

***ALTER TABLE sales NO
INMEMORY;***



Why is In-Memory faster than buffer cache?




SELECT **COL4** FROM MYTABLE;



RESULT



Why is In-Memory faster than buffer cache?



SELECT **COL4** FROM MYTABLE;



RESULT

- In the column store the query behaves very differently
- Data is store is separate physical column structures
- Just go directly to the col4 structure and scan all the entries, one after the other.



Dual Format

